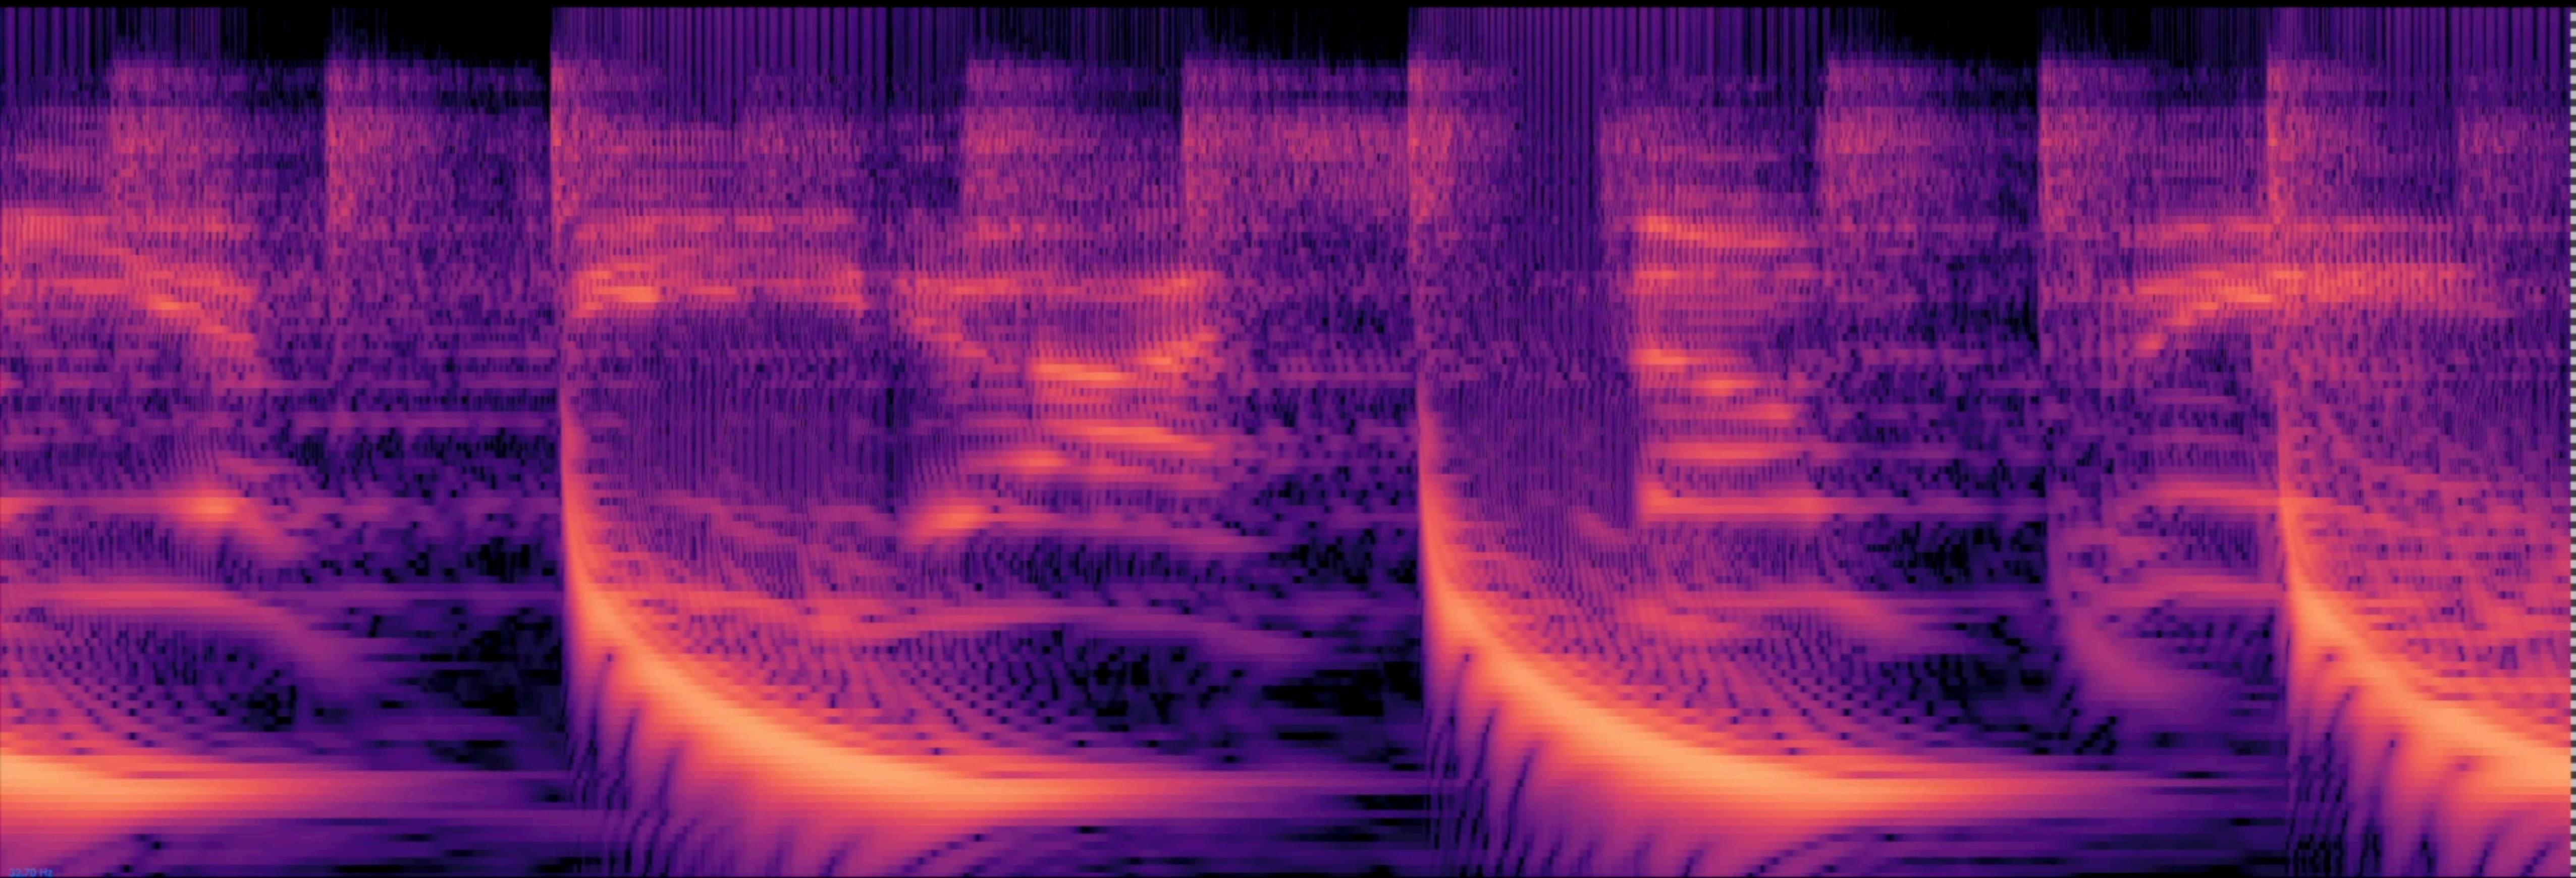


REAL-TIME, LOW LATENCY AND HIGH TEMPORAL RESOLUTION SPECTROGRAMS

ALEXANDRE R.J. FRANCOIS



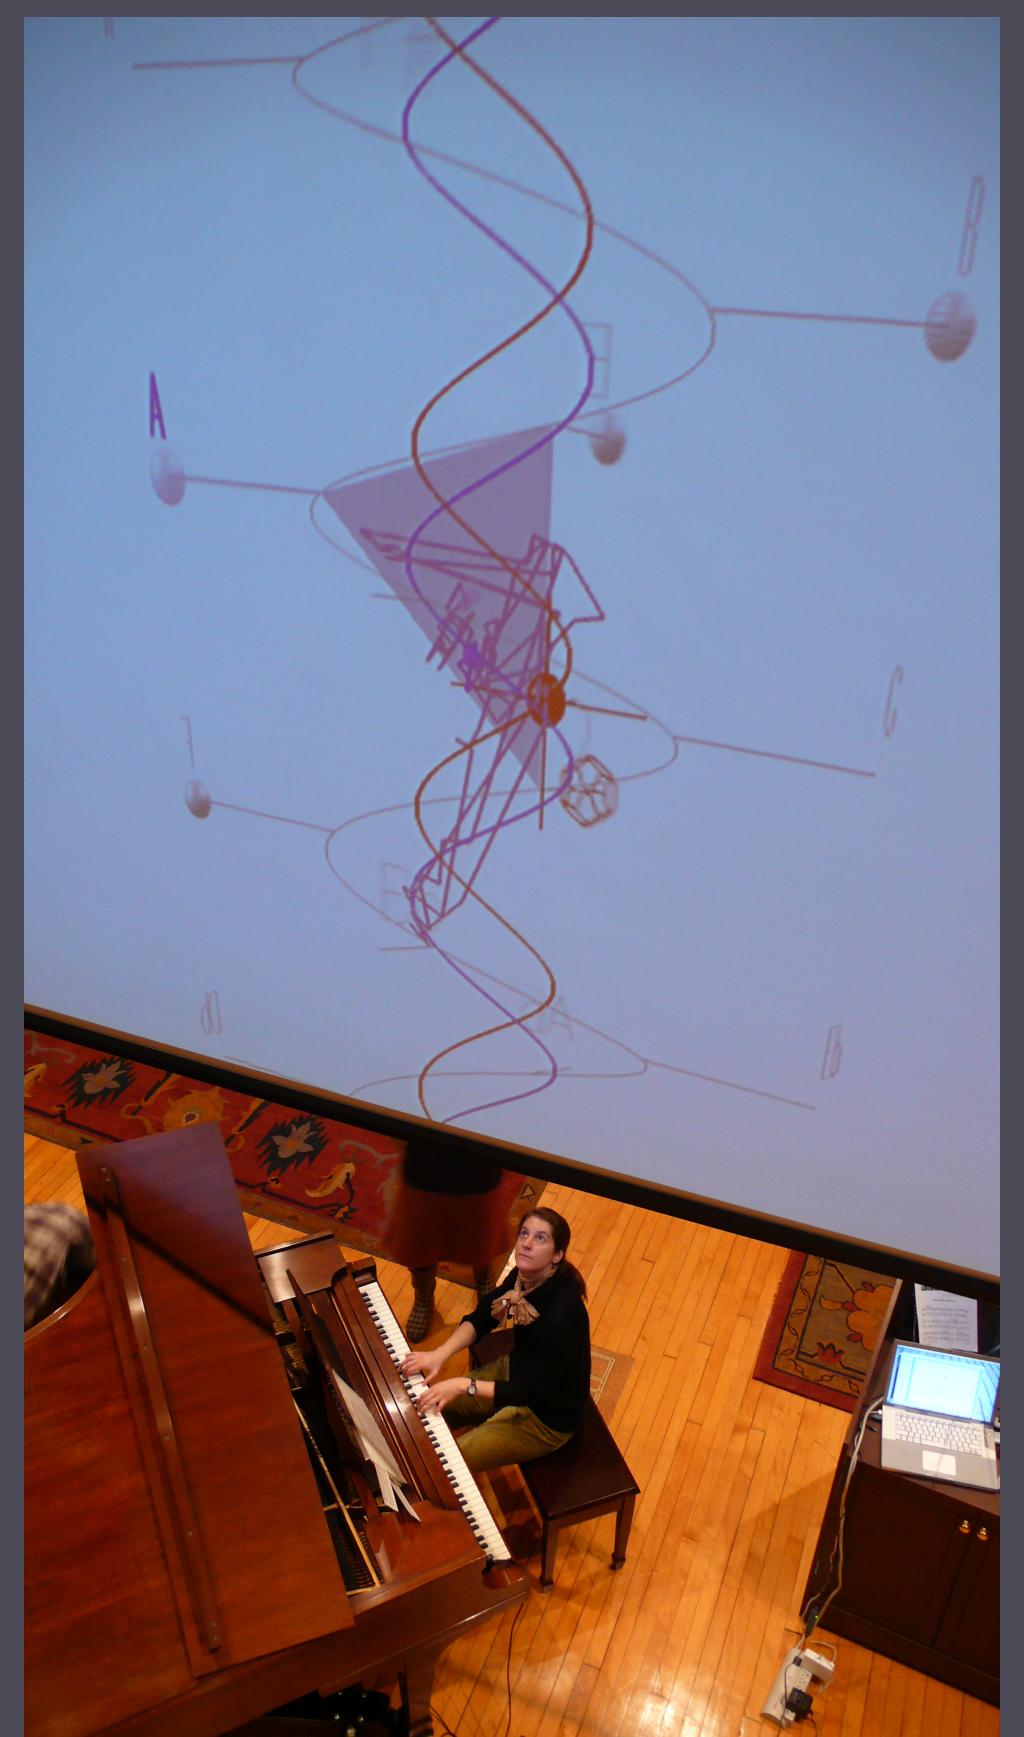
Pump Up The Jam - Technotronic

Alexandre R.J. François

- Computer scientist, software engineer, etc.
- Computer vision, interactive systems, computer graphics, **music computing**
- Academia, startups, larger companies
- Interactive systems, design, user experience (with computers)
 - For **brains**

www.alexandrefrancois.org

MuSA.RT (2003+) with E. Chew



MIMI (2007+) with E. Chew

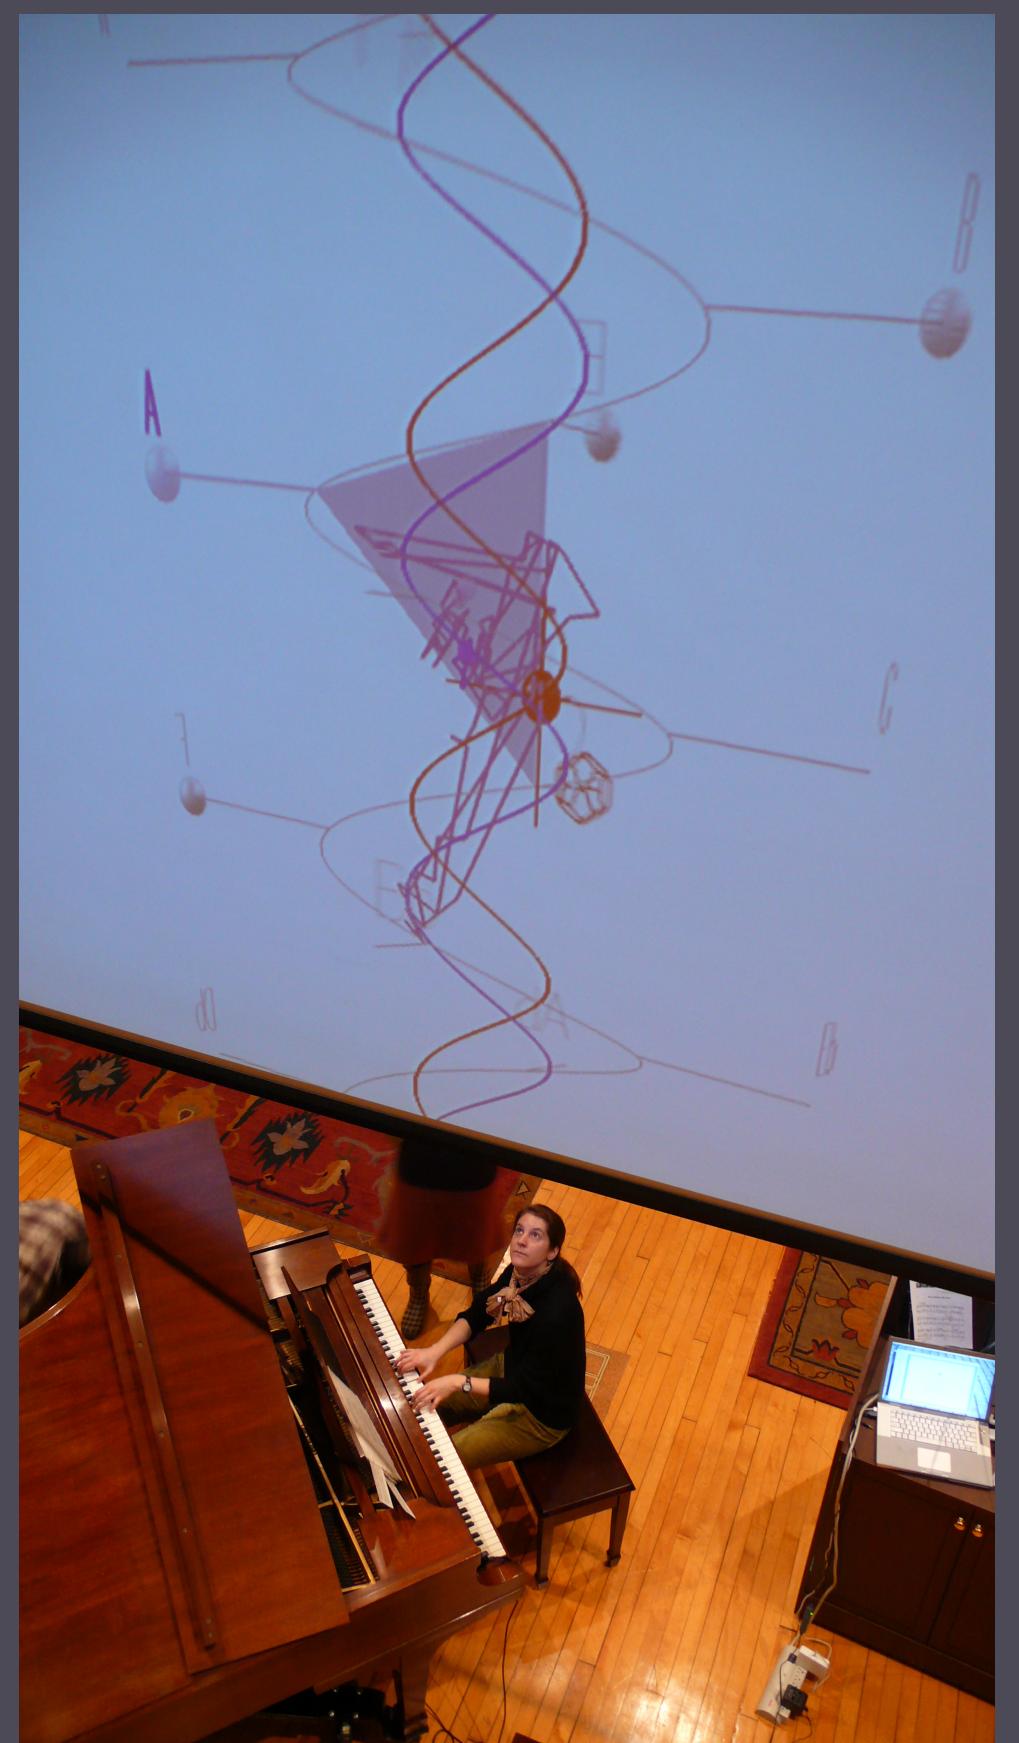


Context and Motivation

Which/how much of the frequencies of interest are present in the input signal?

- Input to interactive music analysis systems for visualization, improvisation, etc.
 - MIDI
 - Perceptually relevant
 - Low latency
 - Computationally efficient

MuSA.RT (2003+) with E. Chew



MIMI (2007+) with E. Chew



The Fast Fourier Transform (FFT)

But... is it really the best solution?

- Fast, but... buffering imposes:
 - Latency
 - Linearly distributed frequency bins
 - Time / frequency resolution trade-off
- Mitigation techniques can be computationally expensive
 - Sliding window
 - Constant-Q transform (log frequency scale)

Buffer size (samples)	Duration (ms)
1	0.023
8	0.181
16	0.363
32	0.726
64	1.451
128	2.903
256	5.805
512	11.610
1024	23.220
2048	46.440
4096	92.880

Sampling rate: 44.1 kHz

A Different Approach: Resonate

- Simple resonator model, assembled into banks
 - Efficient iterative formulation and implementation, no buffering
- Analogous to a bank of non-linear filters
- Similar/related/identical? to Goertzel algorithm, SWIFT
- Not a better FFT - not trying to compute the DFT
- Not aware of this approach being used in a practical way for music/audio
- Reference publication:
Alexandre R.J. François, “Resonate: Efficient Low Latency Spectral Analysis of Audio Signals,” in Proceedings of the 50th Anniversary of the International Computer Music Conference 2025, pp. 251-258, Boston, MA, USA, 8-14 June 2025.

Resonate Properties

- Direct computation of spectrum with arbitrary frequency distribution
- High temporal resolution and low latency: one measurement per input sample
- Resonators independently tuned: frequency and dynamics
- Works across whole frequency range with same efficiency
 - Including low frequencies
 - ...and the resulting data look reasonable...

This Talk

- Goal: share what I have done, no matter how imperfect, so people can use it, improve on it, dream up and build applications aligned with their own interests...
- This is a personal project rooted in my past academic work, developed on week-ends, holidays and other breaks, over a number of years
 - Related open source material online since 2022
 - Does not represent the views of past or present employers

Resources

- *noFFT* open source module: python and C++ implementations, Jupyter notebooks with code used to generate the figures in the paper (and more...)
- Oscillators open source Swift package: reference implementations in Swift and C++
- **Oscillators app for iOS/iPadOS/MacOS (Apple Silicon, designed for iPad)**
- Resonate Youtube playlist

www.alexandrefrancois.org/Resonate

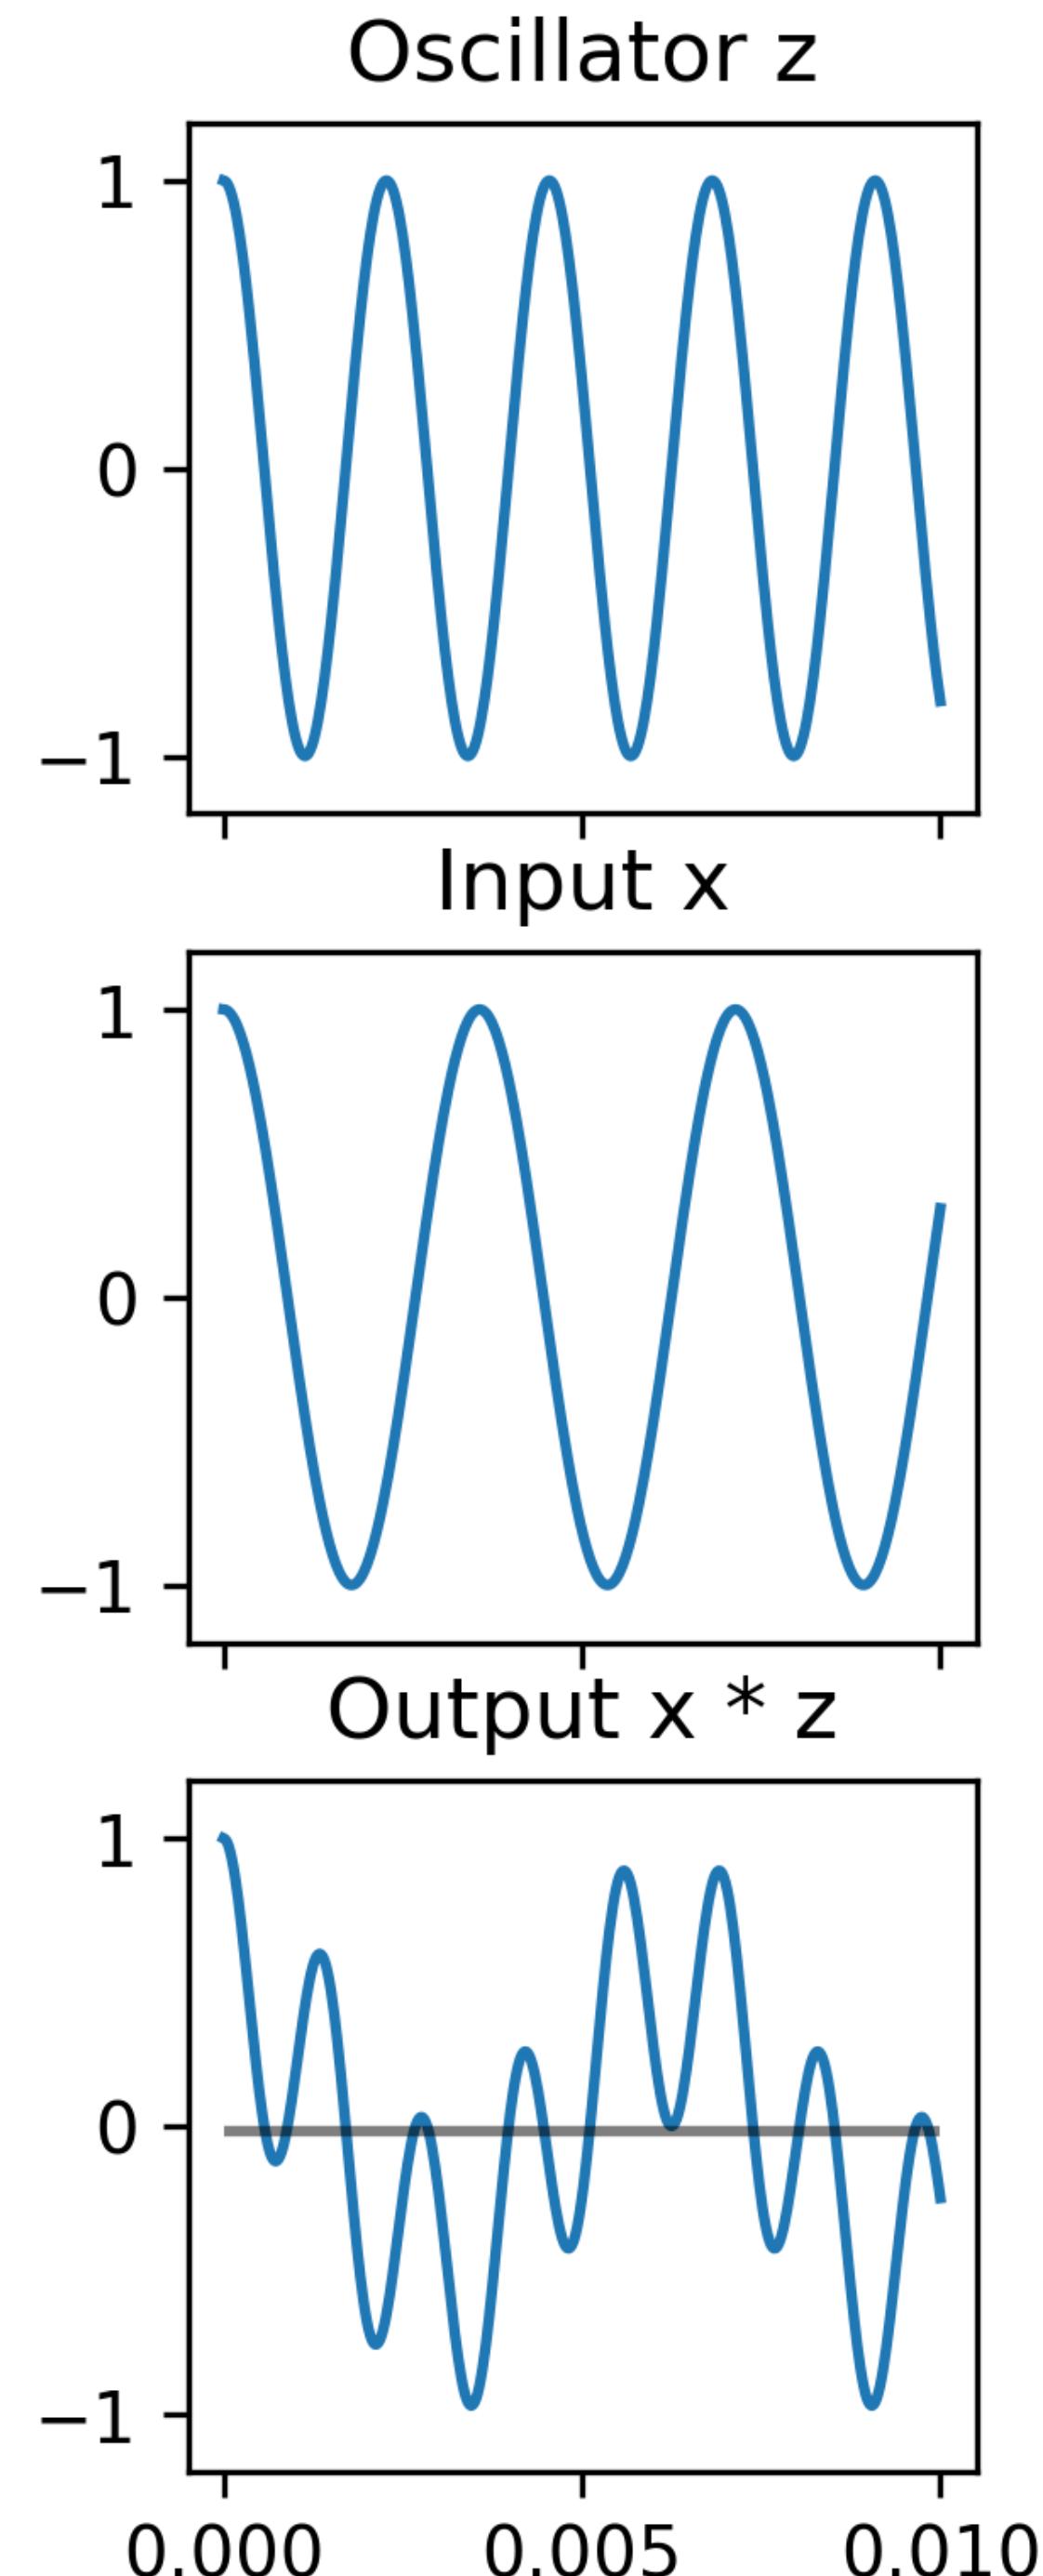
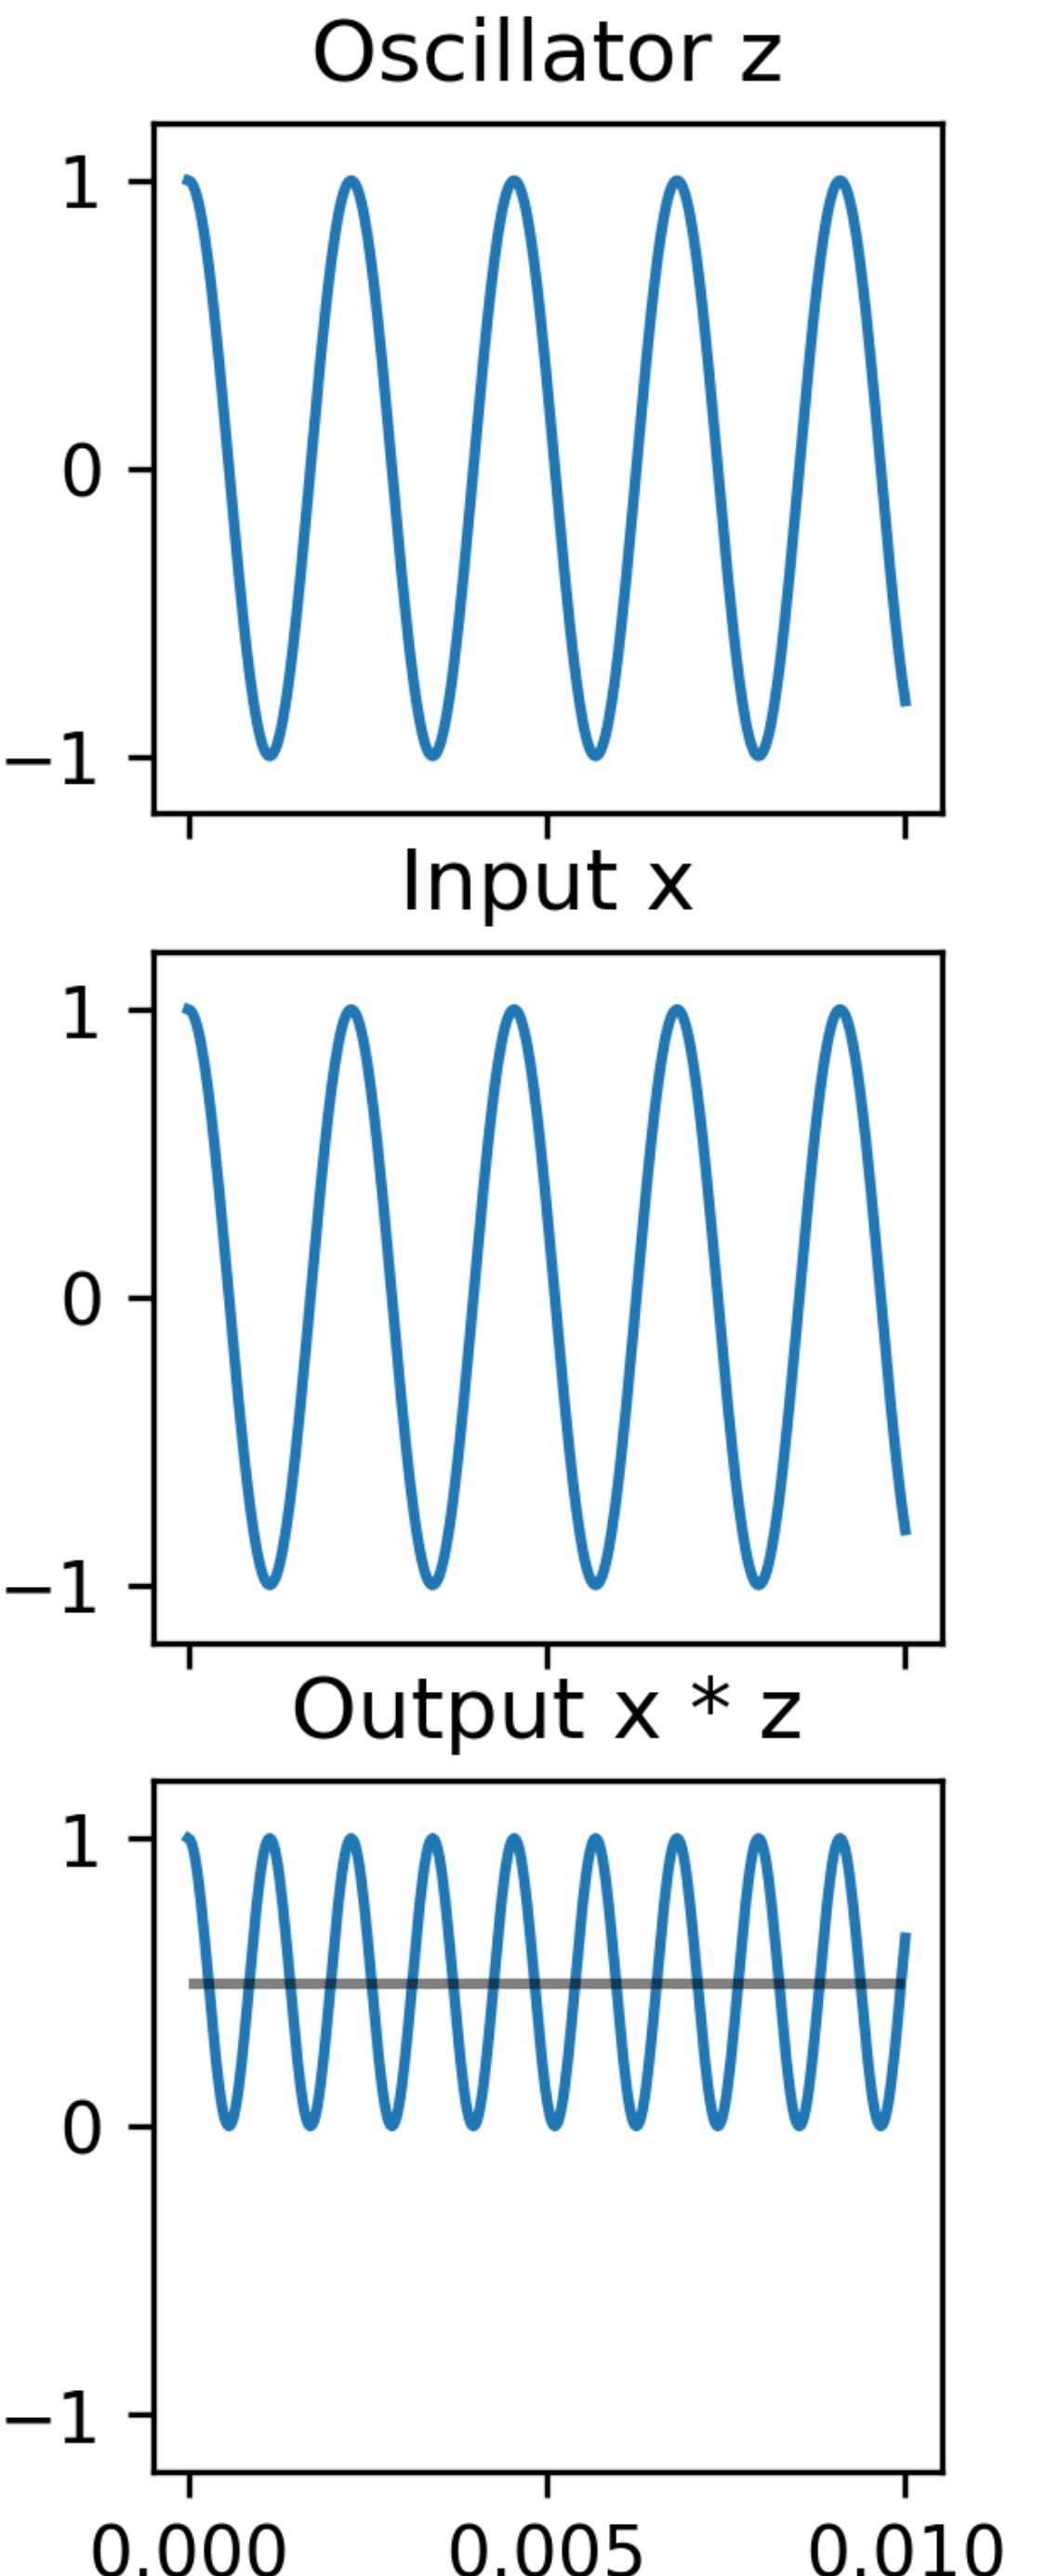


Outline

- Intro
- **Single resonator**
- Resonator banks
- Apps and applications

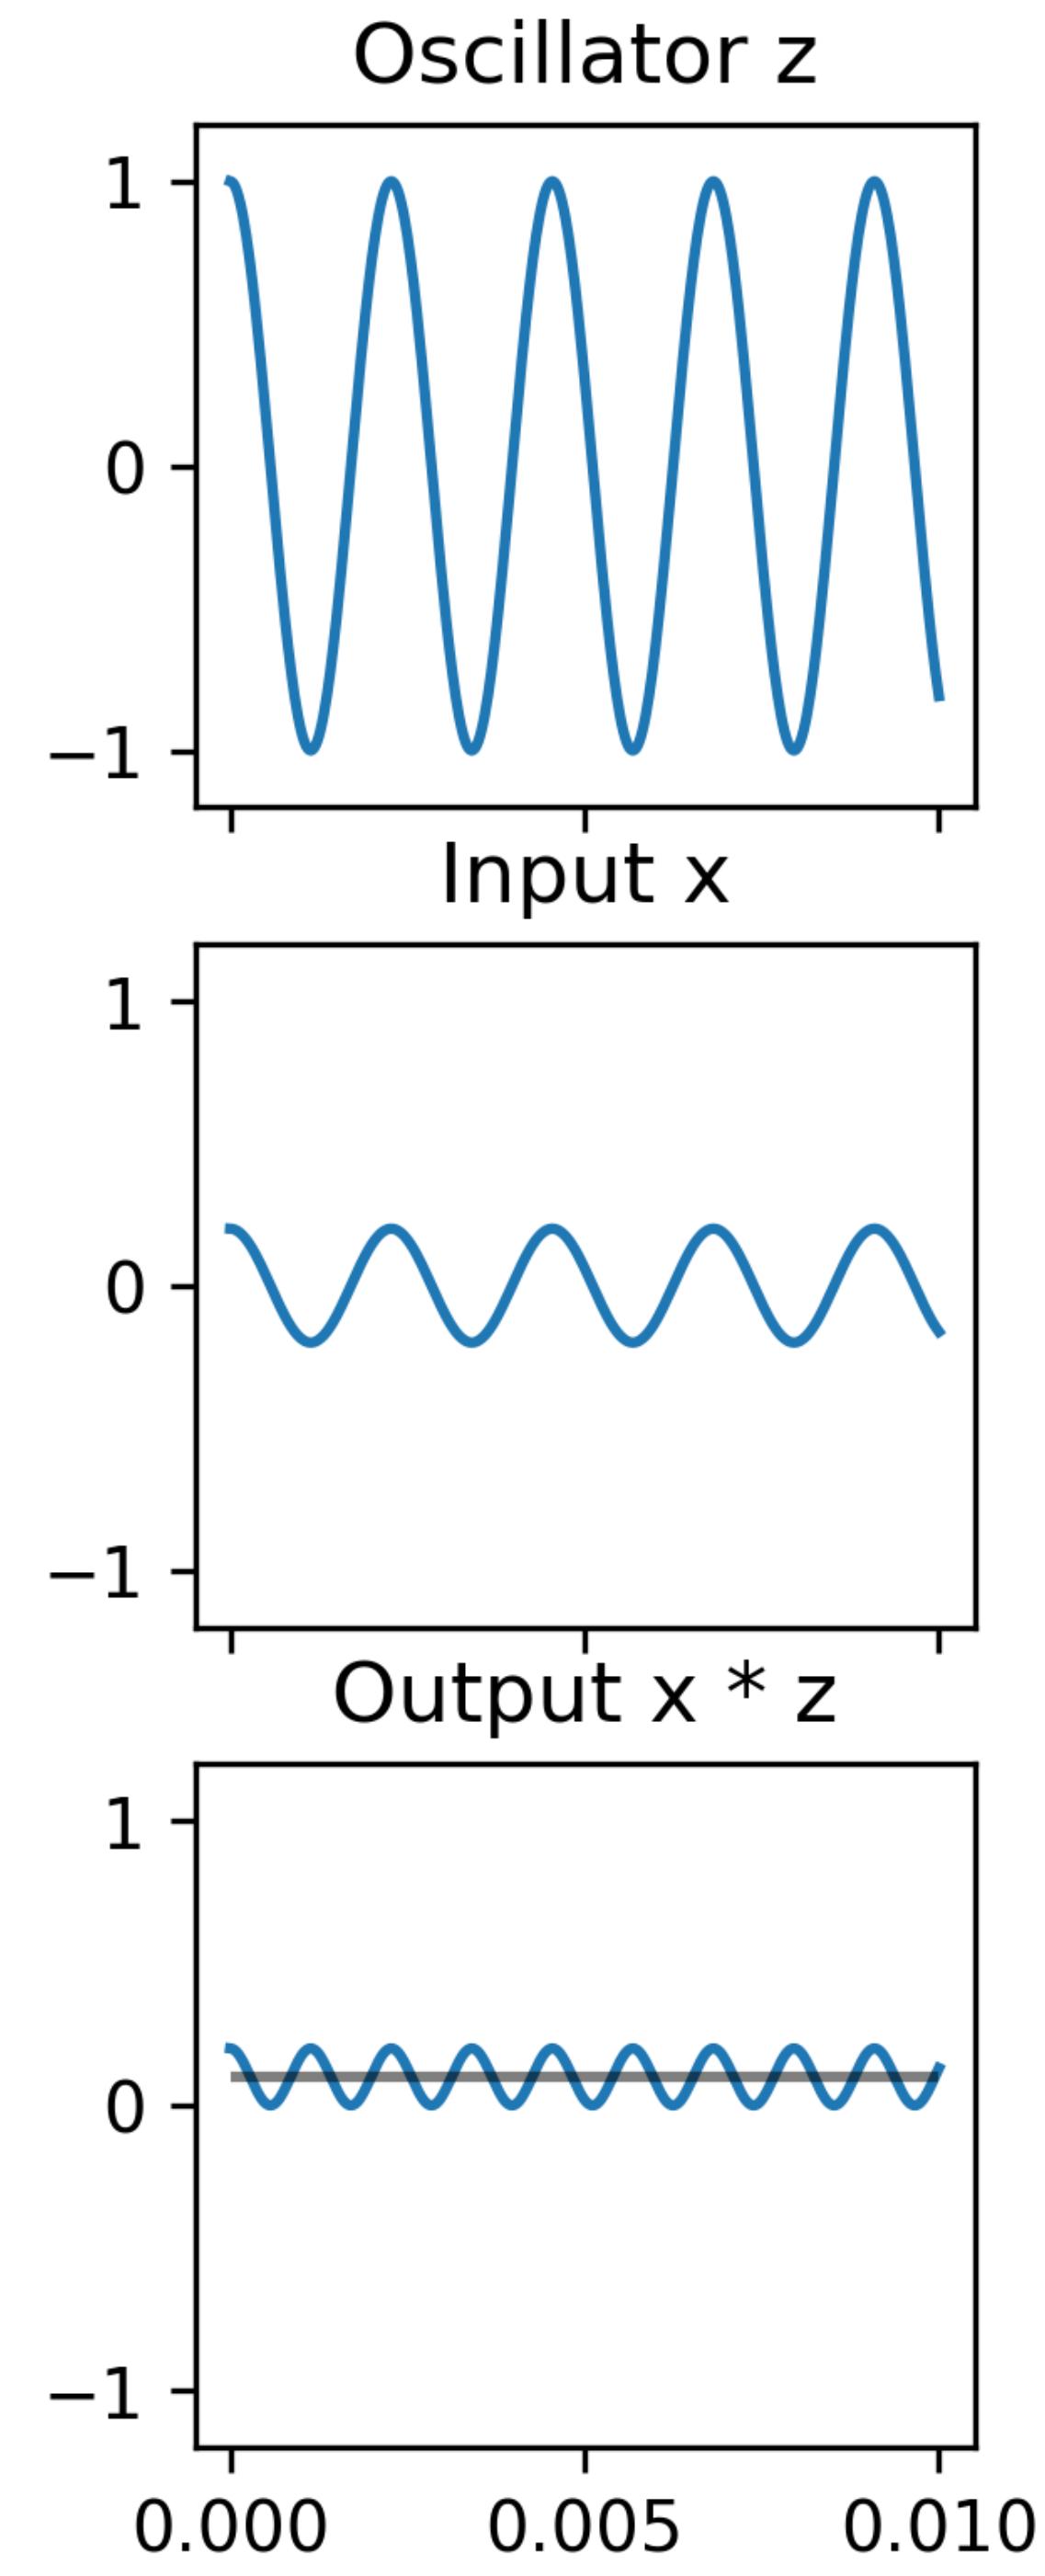
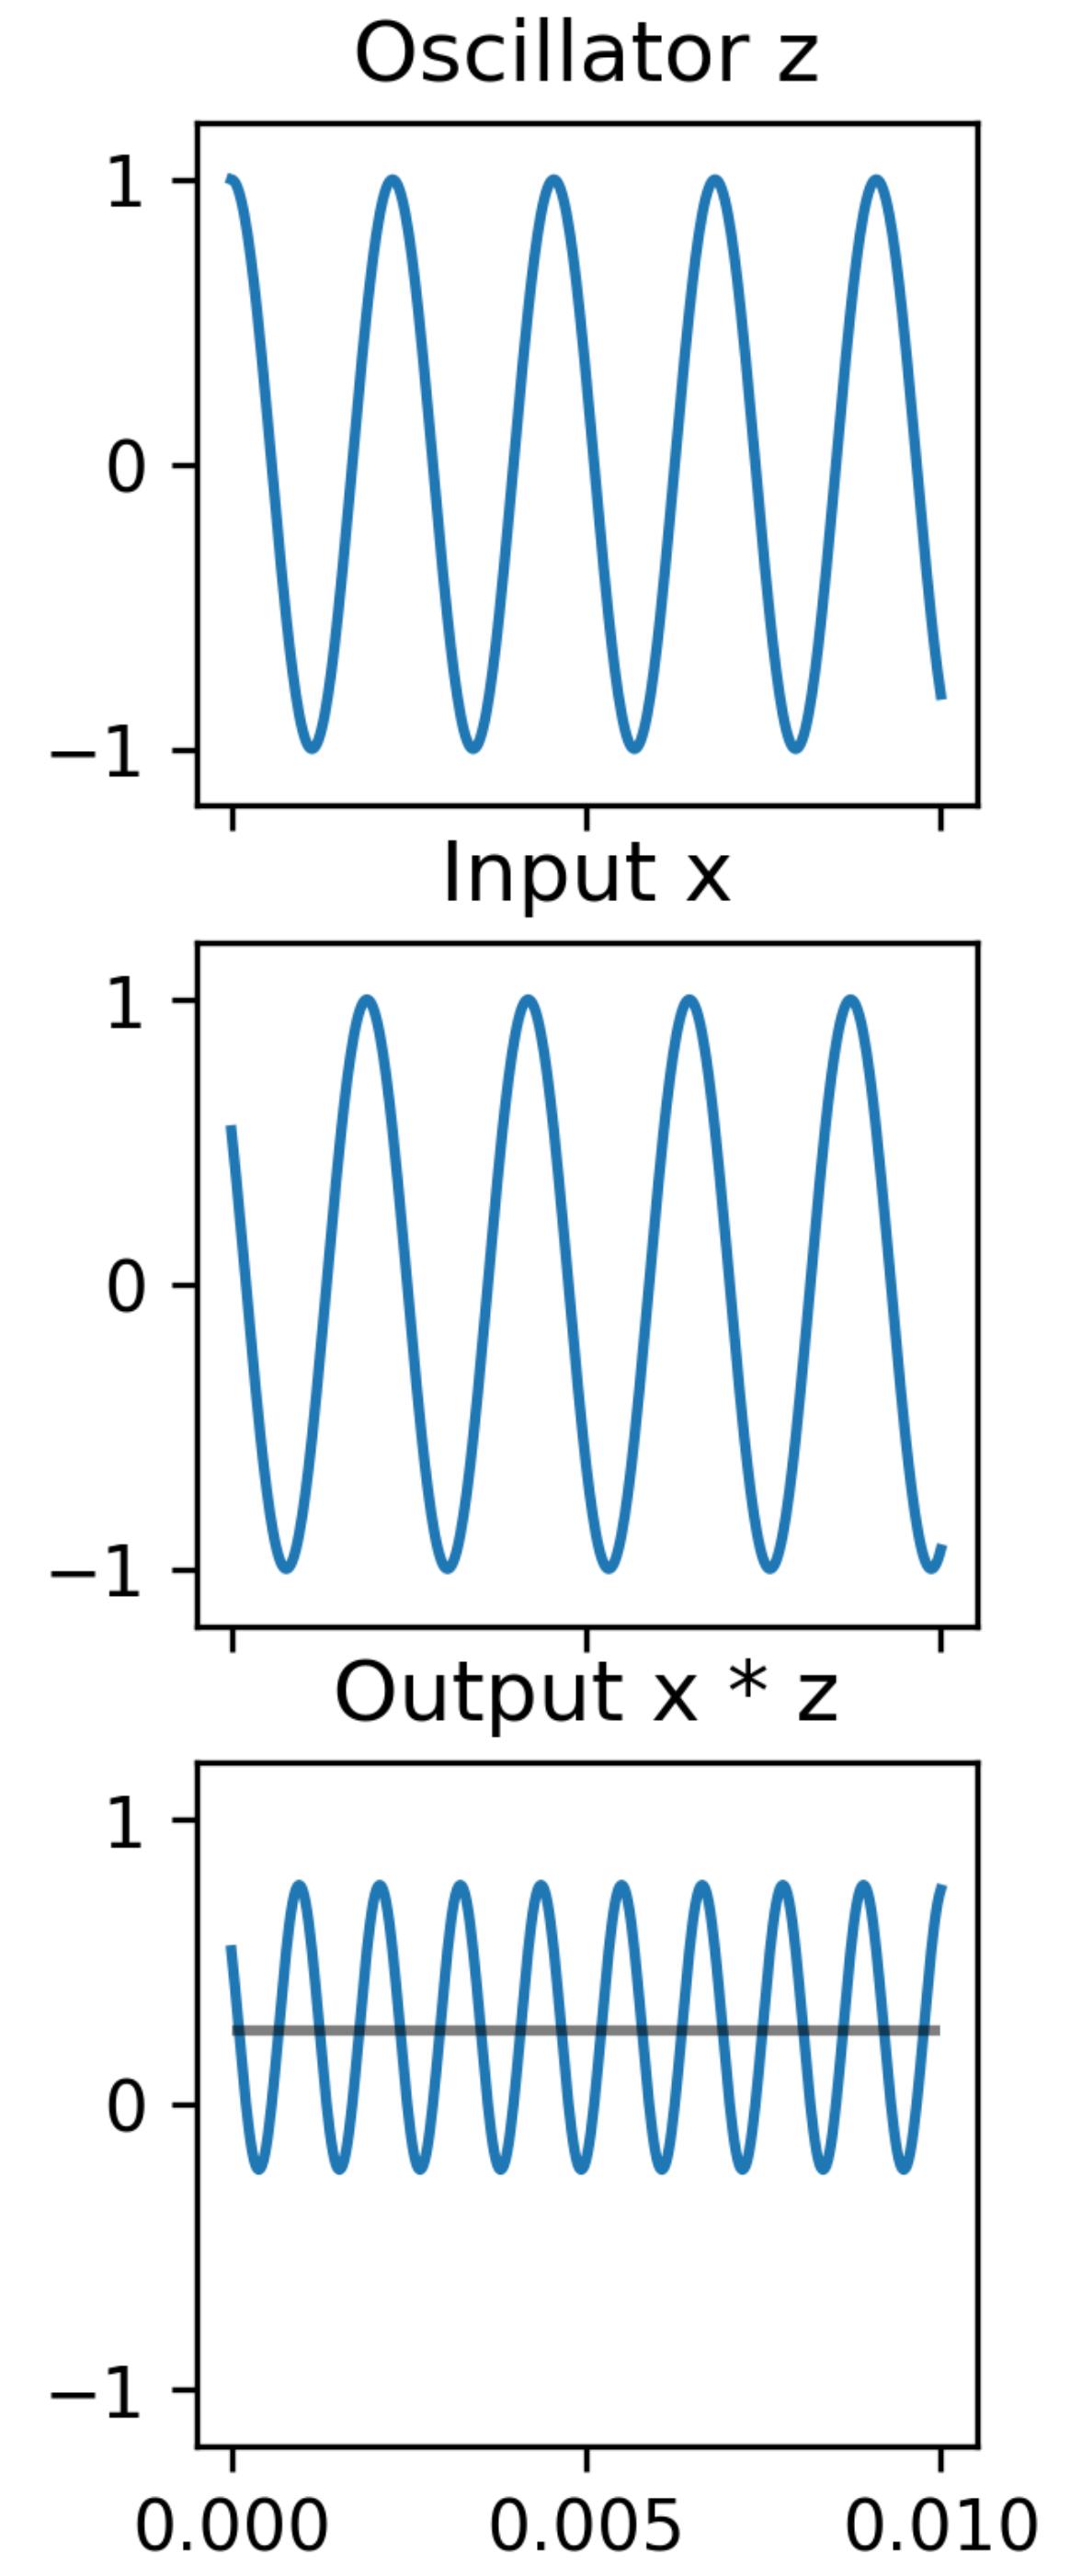
Resonator

- A resonator is a device or system that exhibits resonance, meaning it oscillates with greater amplitude at certain frequencies than others
- Simple example:
 - Oscillator with sine wave z at known frequency (amplitude 1)
 - Input sine wave x (amplitude 1)
 - Output: product of input and oscillator, averaged over time
 - Intuitively, output will be maximal when input and resonator have same frequency... and phase



Phase and Amplitude

- Given resonant frequency
- Input signal of same frequency
- Two parameters to compute:
 - Phase
 - Amplitude



Phase and Amplitude

- 2 degrees of freedom
- Need 2 measurements per input sample
- Use cos and sin
 - $\cos^2(x) + \sin^2(x) = 1$

```
// input signal X
// oscillator waveforms Zc (cosine)
// and Zs (sine)
// for each input sample x
c = x * Zc
s = x * Zs
power = c*c + s*s
amplitude = sqrt(c*c + s*s)
phase = atan2(s, c)
```

Phasor

Iterative implementation of sine oscillator

```

class Phasor {
    var frequency: Float
    var sampleRate: Float
    // Phasor variables
    // Phasor: Z = Zc + j Zs
    // Multiplier: W = Wc + j Ws

    var Zc : Float = 1.0
    var Zs : Float = 0.0
    var Wc : Float = 0.0
    var Ws : Float = 0.0
    // pre-computed Wc + Ws
    var Wcps : Float = 0.0
    [...]
}

    // pre-compute Wc + Ws
    let omega = twoPi * frequency / sampleRate
    Wc = cos(omega)
    Ws = sin(omega)
    Wcps = Wc + Ws

    // Compute next value of the phasor
    // Z <- Z * W
    // complex multiplication
    // with 3 real multiplications
    let ac = Wc*Zc
    let bd = Ws*Zs
    let abcd = Wcps * (Zc+Zs)
    Zc = ac - bd
    Zs = abcd - ac - bd

```

Phasor

Iterative implementation of sine oscillator

```

class Phasor {
    var frequency: Float
    var sampleRate: Float
    // Phasor variables
    // Phasor: Z = Zc + j Zs
    // Multiplier: W = Wc + j Ws
    var Zc : Float = 1.0
    var Zs : Float = 0.0
    var Wc : Float = 0.0
    var Ws : Float = 0.0
    // pre-computed Wc + Ws
    var Wcps : Float = 0.0
    [...]
}

    // Apply re-normalization correction
    // to compensate for numerical drift
    // Use Taylor expansion of 1/sqrt(x)
    // around 1 to reduce computational cost
    // This can be applied every n? samples
    let k = (3.0 - Zc*Zc - Zs*Zs) / 2.0
    Zc *= k
    Zs *= k
}

```

Exponentially Weighted Moving Average aka Low-Pass Filter

```
// input signal X  
  
// smoothed signal xx  
// update at each sample x  
  
xx = (1-alpha) * xx + alpha * x  
  
// parameter alpha relates to time constant:  
// time for system response to a unit step function  
// to reach ~62% of the original signal  
  
sampleDuration = 1.0 / sampleRate  
  
tau = -sampleDuration / log(1.0 - alpha)
```

Resonator

EWMA

```
// Update with sample x
class Resonator : Phasor {
    // complex: r = c + j s
    var c: Float = 0.0
    var s: Float = 0.0
    ...
}
```

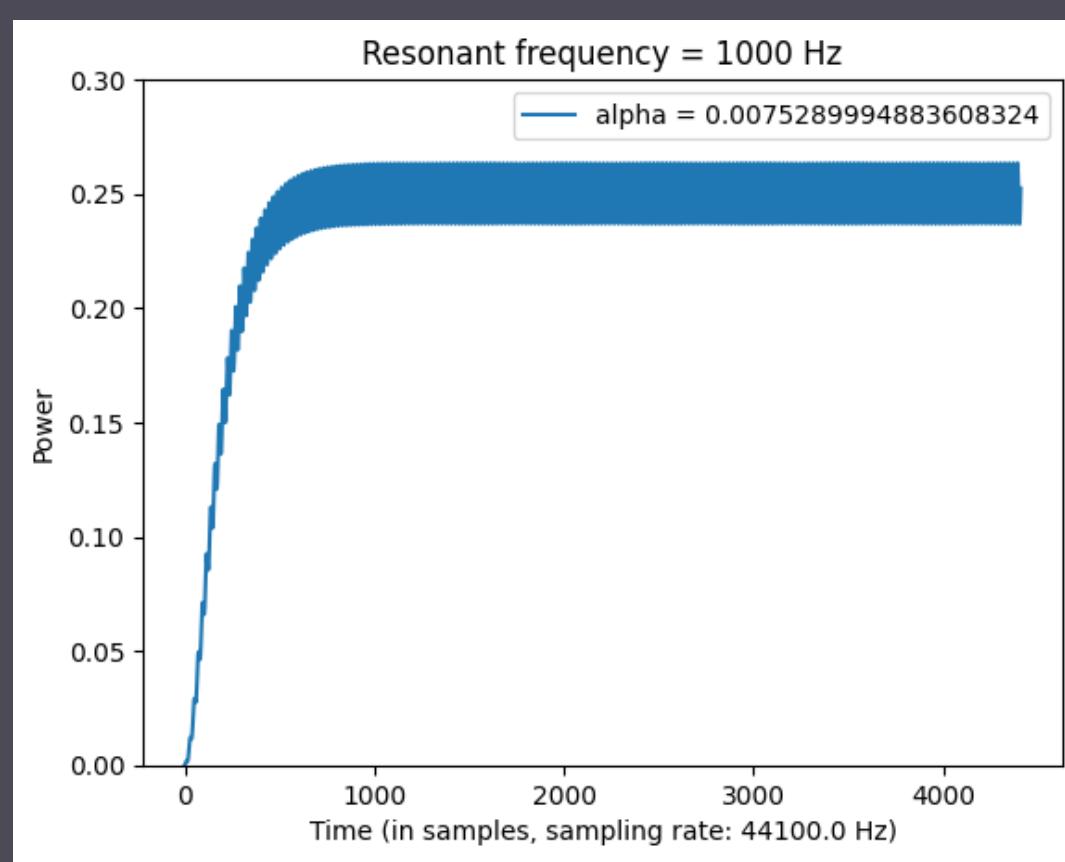
$$\begin{aligned}c &= (1-\alpha) * c + \alpha * x * Z_c \\s &= (1-\alpha) * s + \alpha * x * Z_s\end{aligned}$$

Resonator

EWMA

```
class Resonator : Phasor {
    // complex: r = c + j s
    var c: Float = 0.0
    var s: Float = 0.0

    [...]
}
```



```
// Update with sample
alphaSample = alpha * sample
omAlpha = 1.0 - alpha
c = omAlpha * c + alphaSample * Zc
s = omAlpha * s + alphaSample * Zs
```

```
// Compute next value of the phasor
// Z <- Z * W
// complex multiplication
// with 3 real multiplications

ac = Wc * Zc
bd = Ws * Zs
abcd = (Wcps) * (Zc+Zs)
Zc = ac - bd
Zs = abcd - ac - bd
```

Resonator

Smooth out output oscillations

```

class Resonator : Phasor {
    // complex: r = c + j s
    var c: Float = 0.0
    var s: Float = 0.0

    // Smoothed resonator output
    var cc: Float = 0.0
    var ss: Float = 0.0

    [...]
}

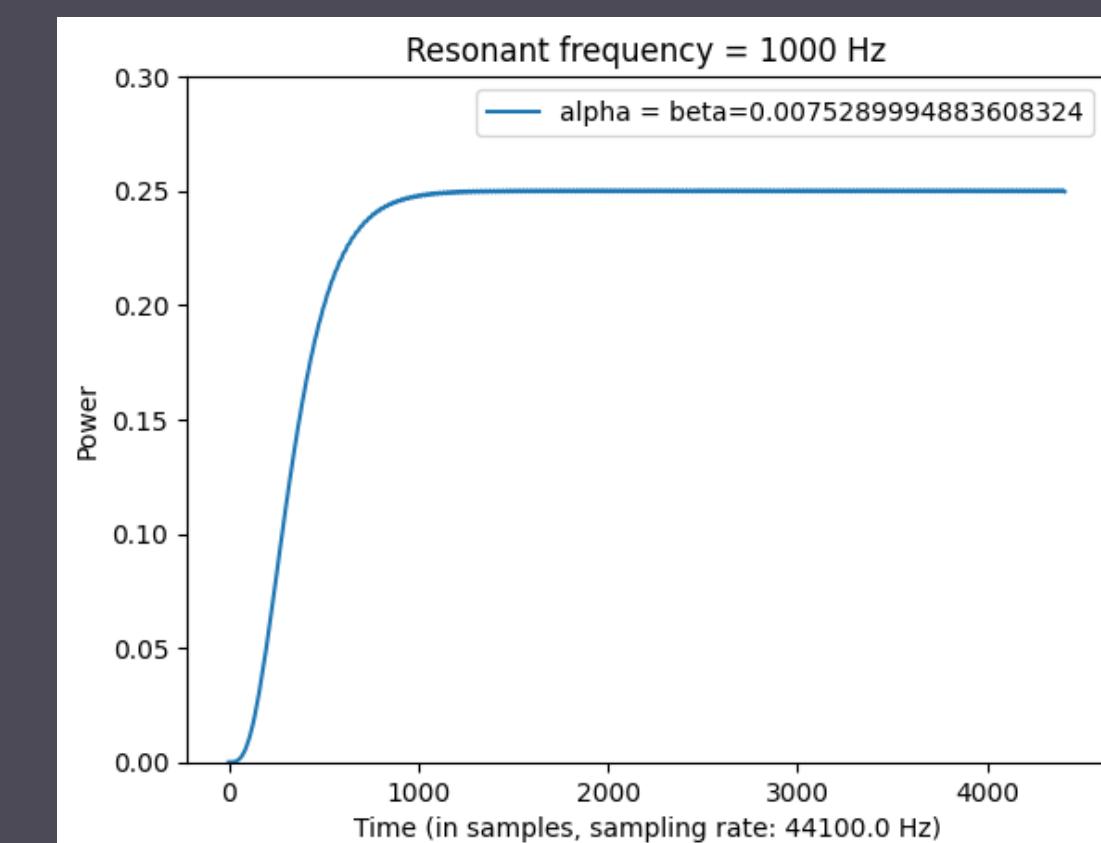
```

```

    // Update with sample
    alphaSample = alpha * sample
    omAlpha = 1.0 - alpha
    c = omAlpha * c + alphaSample * Zc
    s = omAlpha * s + alphaSample * Zs

    omBeta = 1.0 - beta
    cc = omBeta * cc + beta * c
    ss = omBeta * ss + beta * s

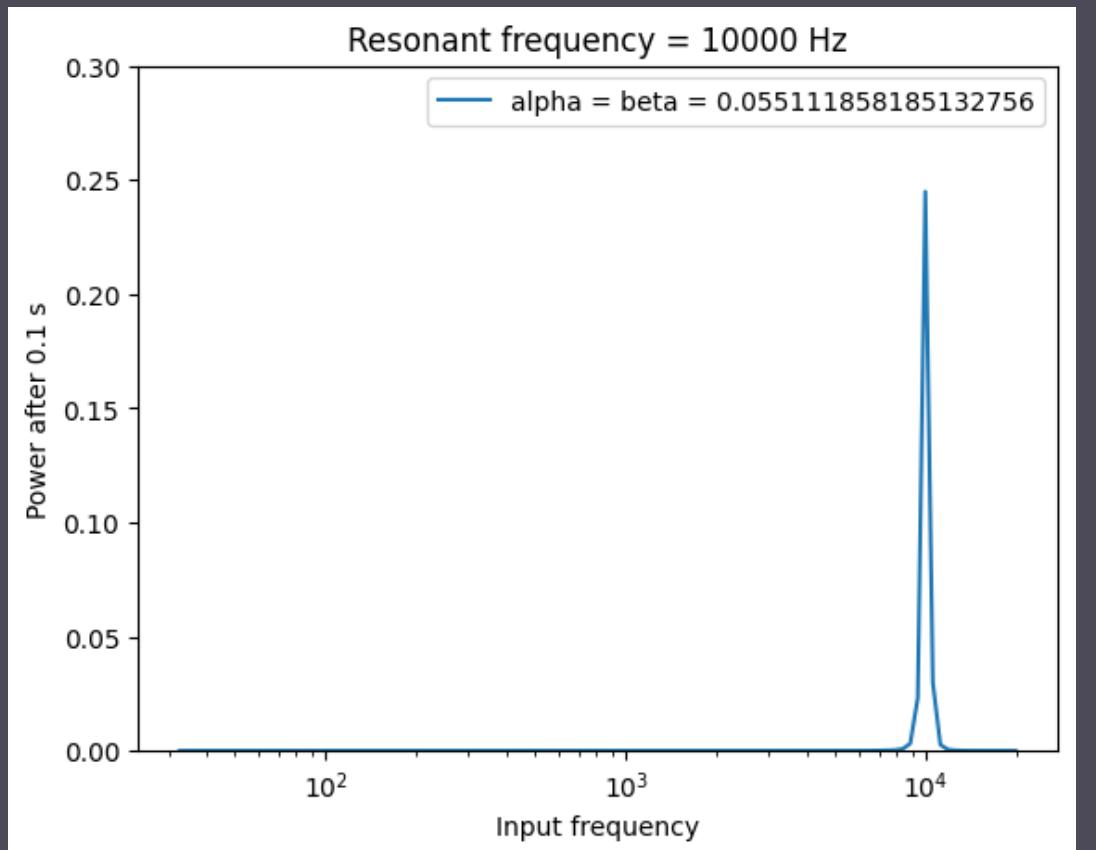
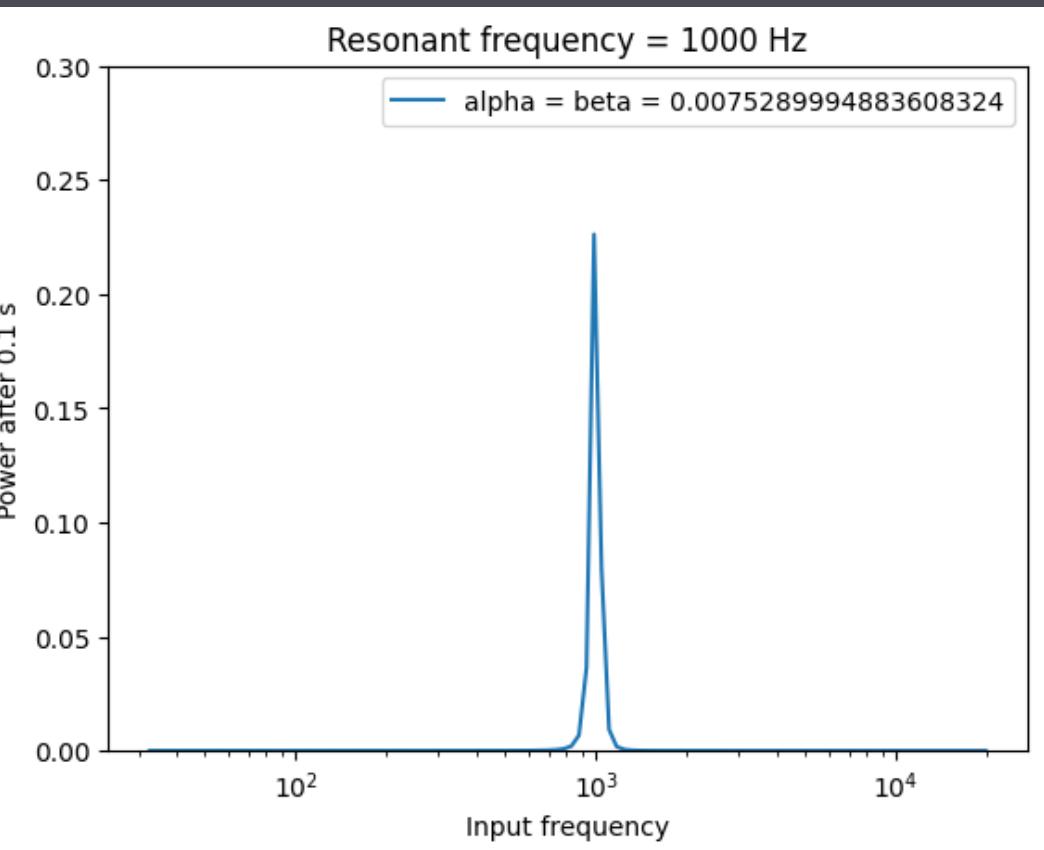
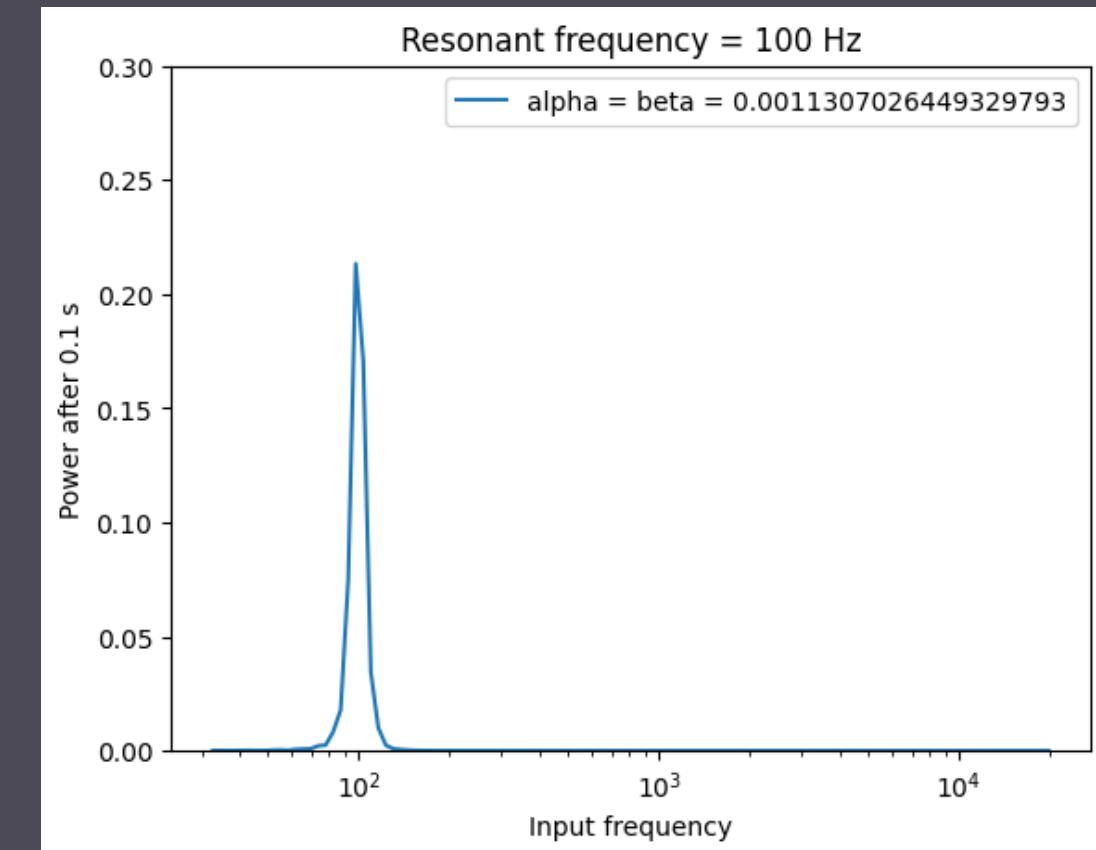
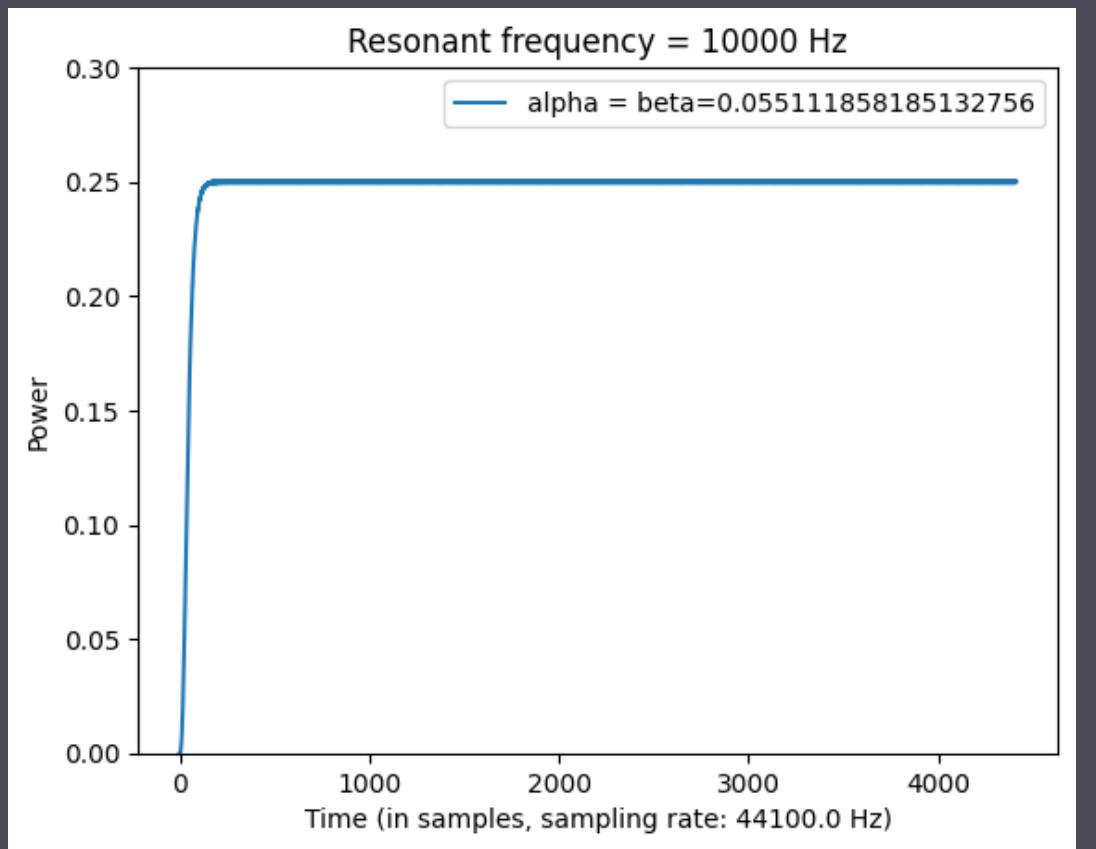
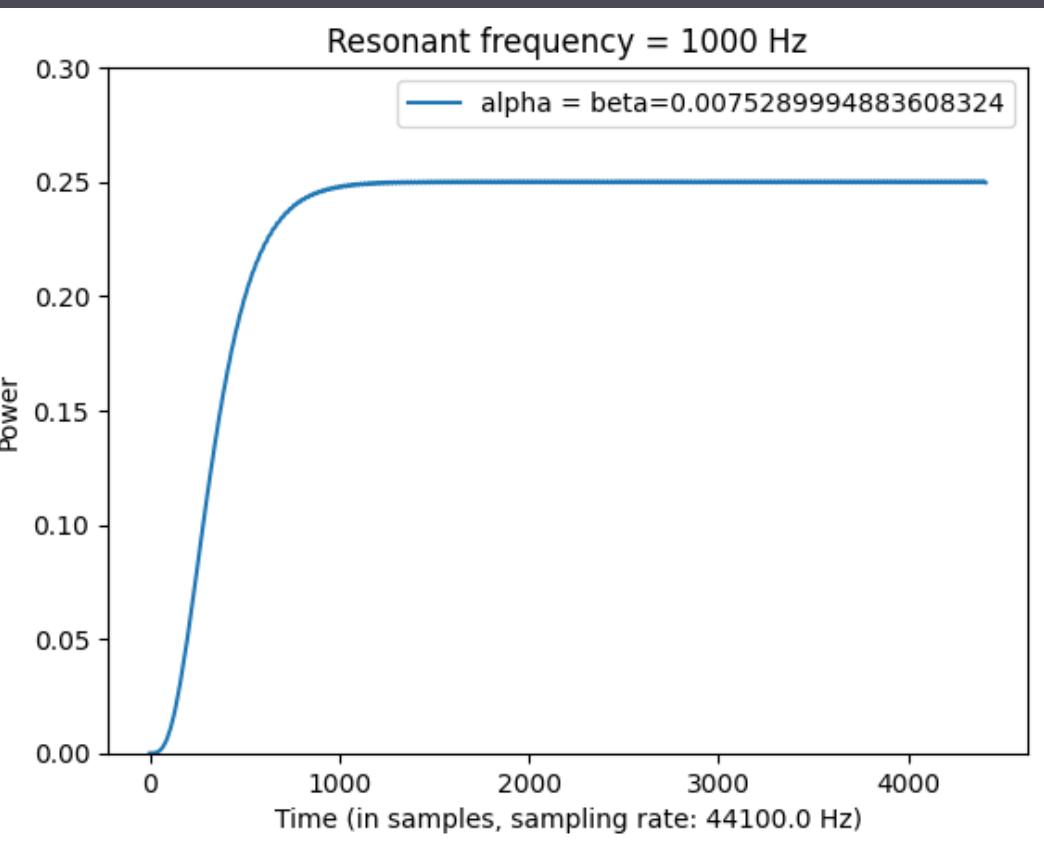
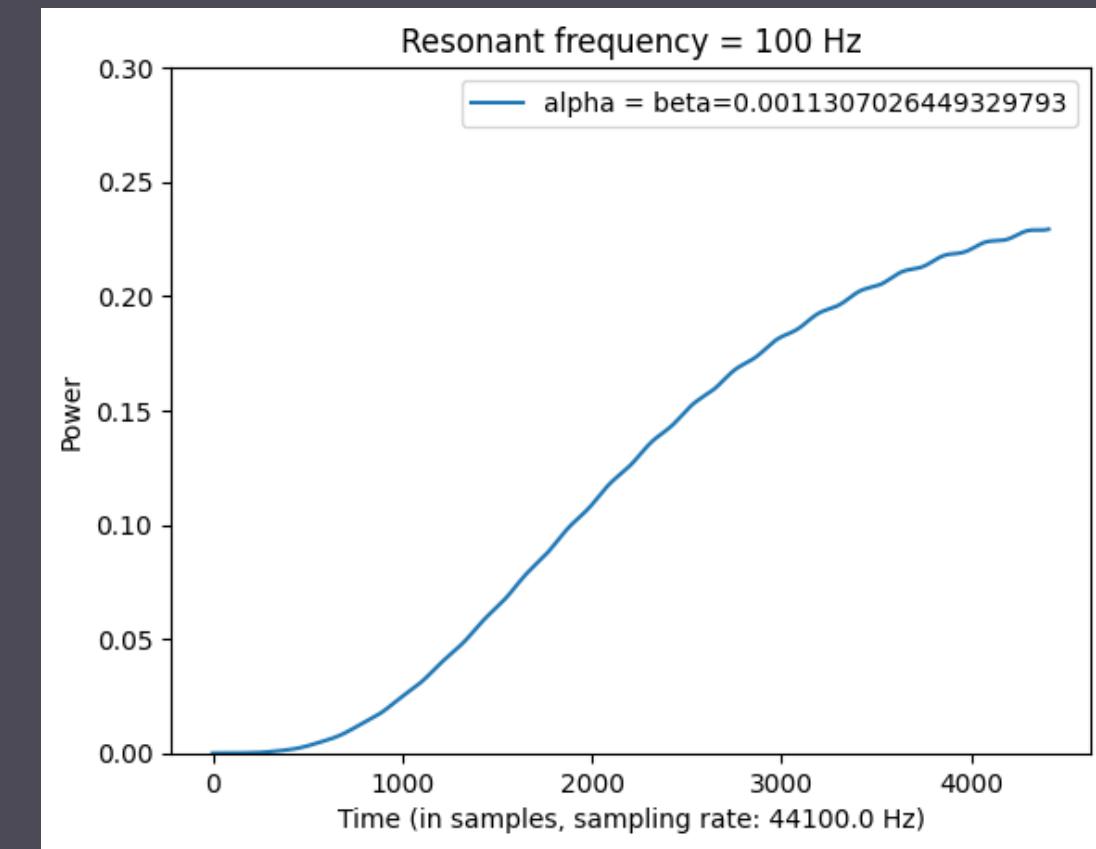
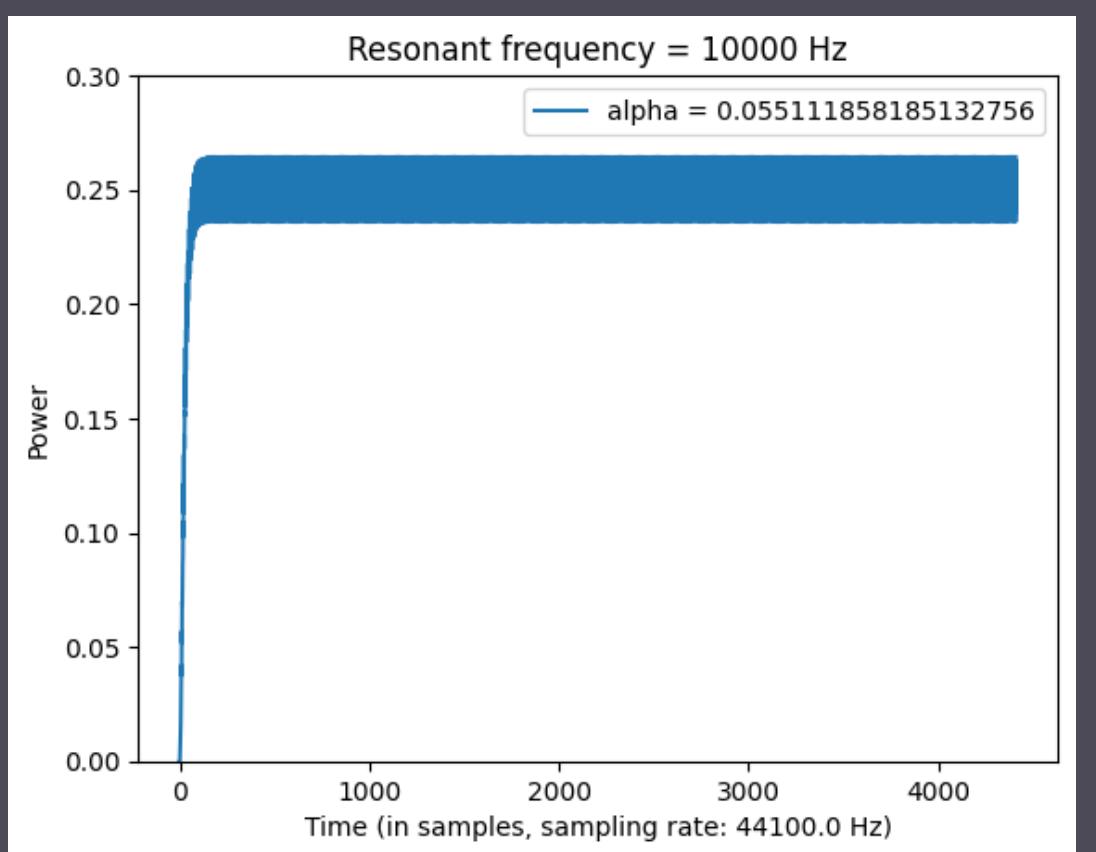
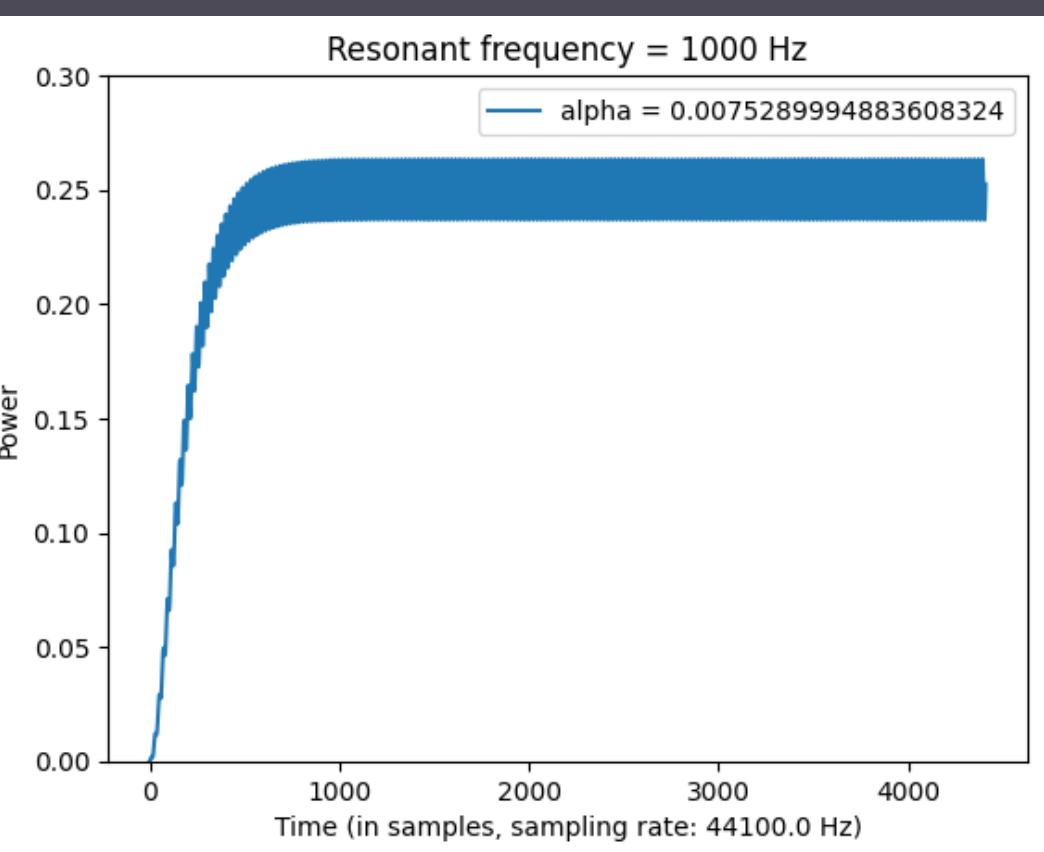
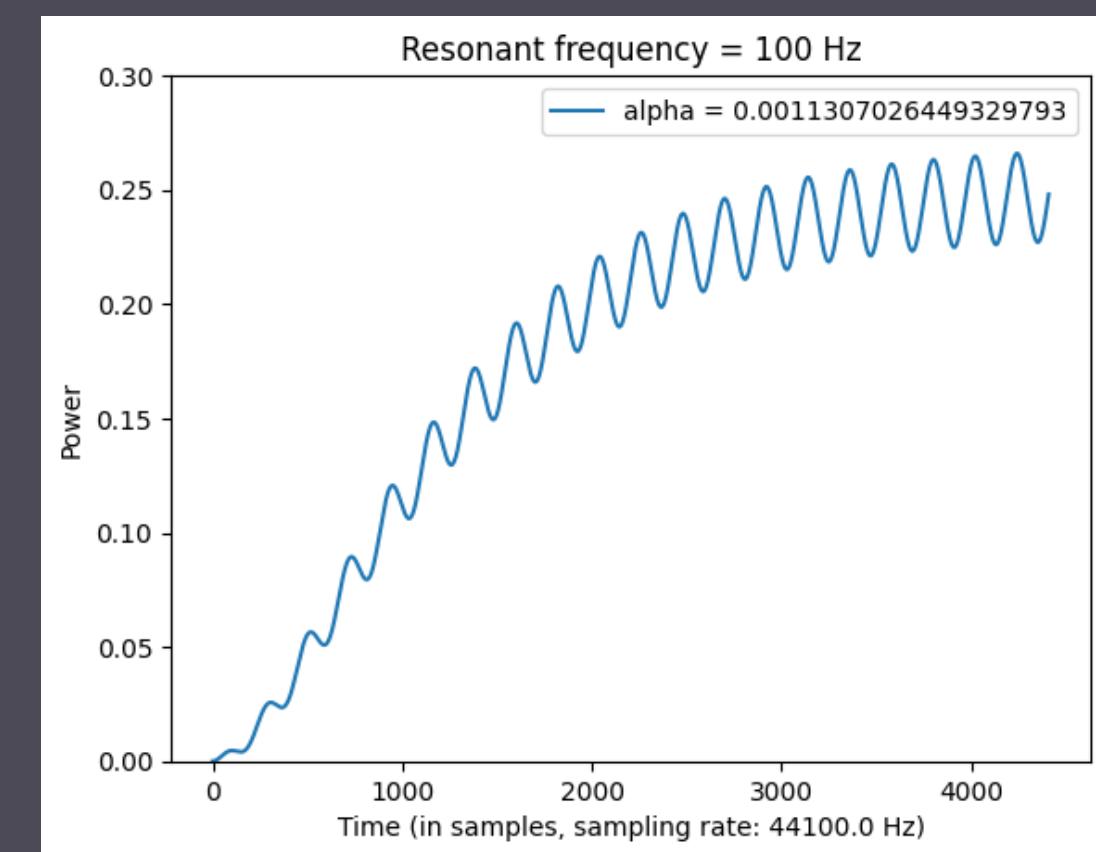
```



Parameter: Time Constant

- Single parameter alpha, set heuristically
- Relate to resonator frequency
- Intuitively, lower frequencies require more time to converge
 - Time constant inversely proportional to frequency

```
tau = k * log10(1+frequency) / frequency  
  
// parameter alpha heuristic  
alpha = 1 - exp(-frequency / (sampleRate * k * log10(1+frequency)))
```

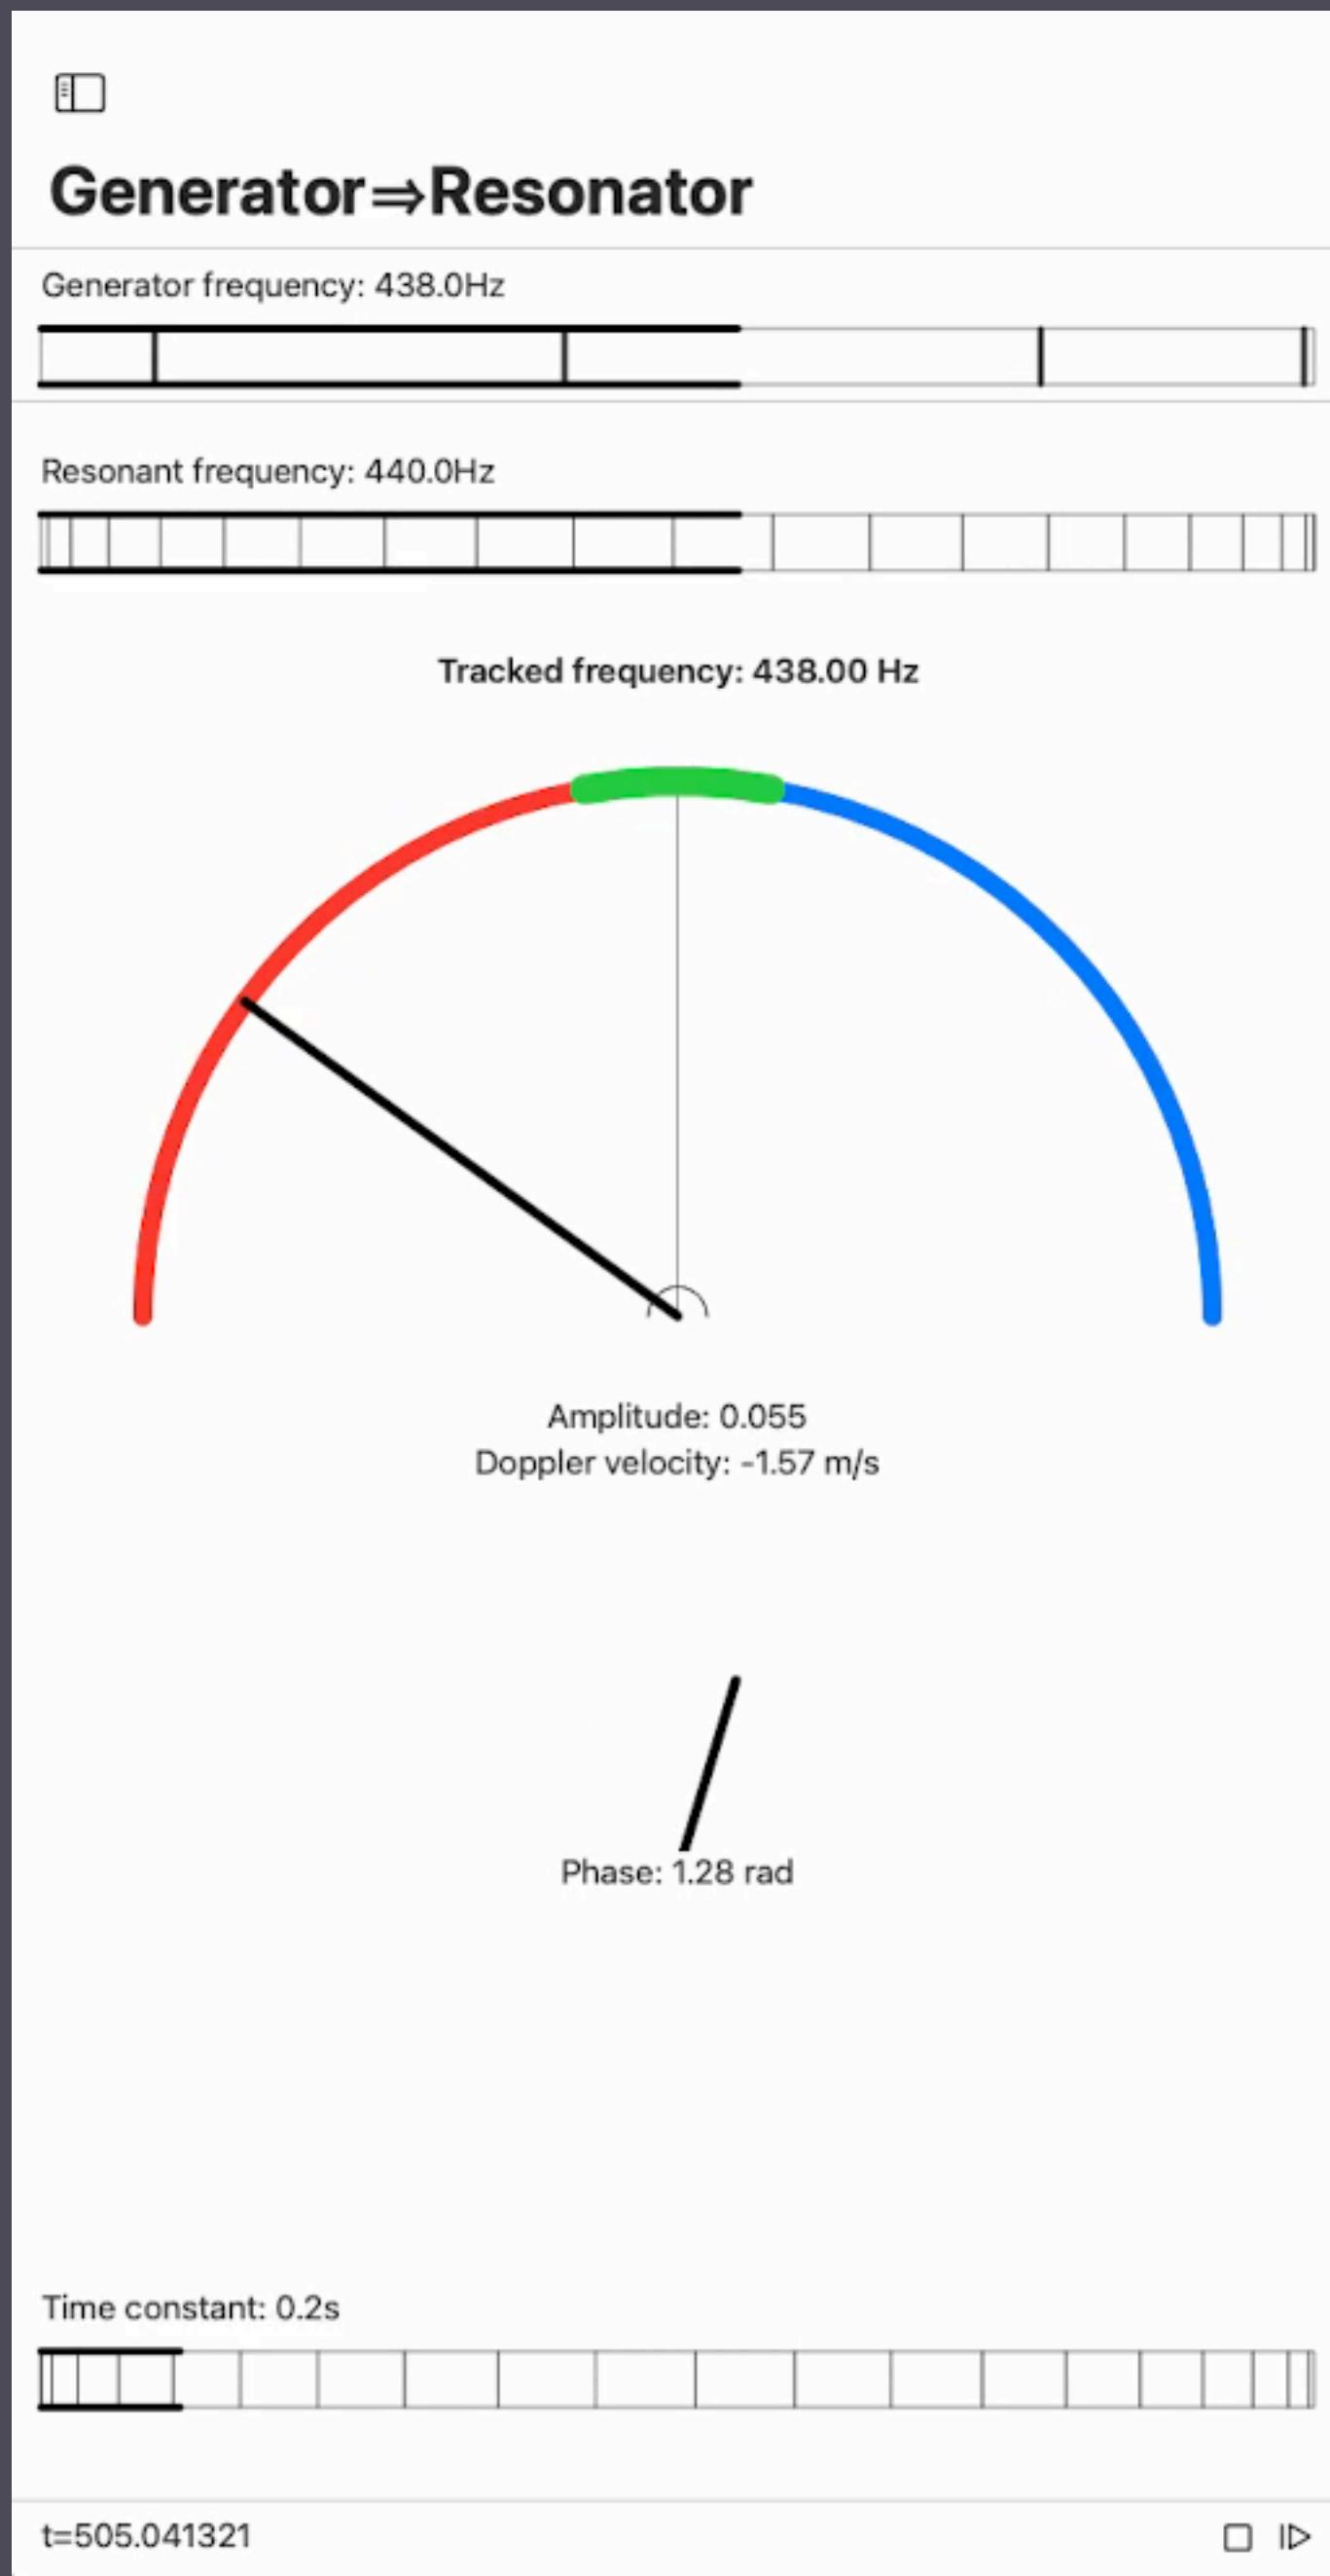


Computational Cost

- Memory: 6 Floats
- Computations per sample update:
 - a handful of multiplications and additions
 - Some occasional overhead to compensate for numerical approximations
 - Complexity is constant
- Overall complexity is linear in the number of samples processed

Some Fun Applications

- Frequency tracking
 - From phase drift (aperture effect)
- Doppler velocity
 - From measured frequency shift
- Stereo directional estimation
 - From stereo phase difference

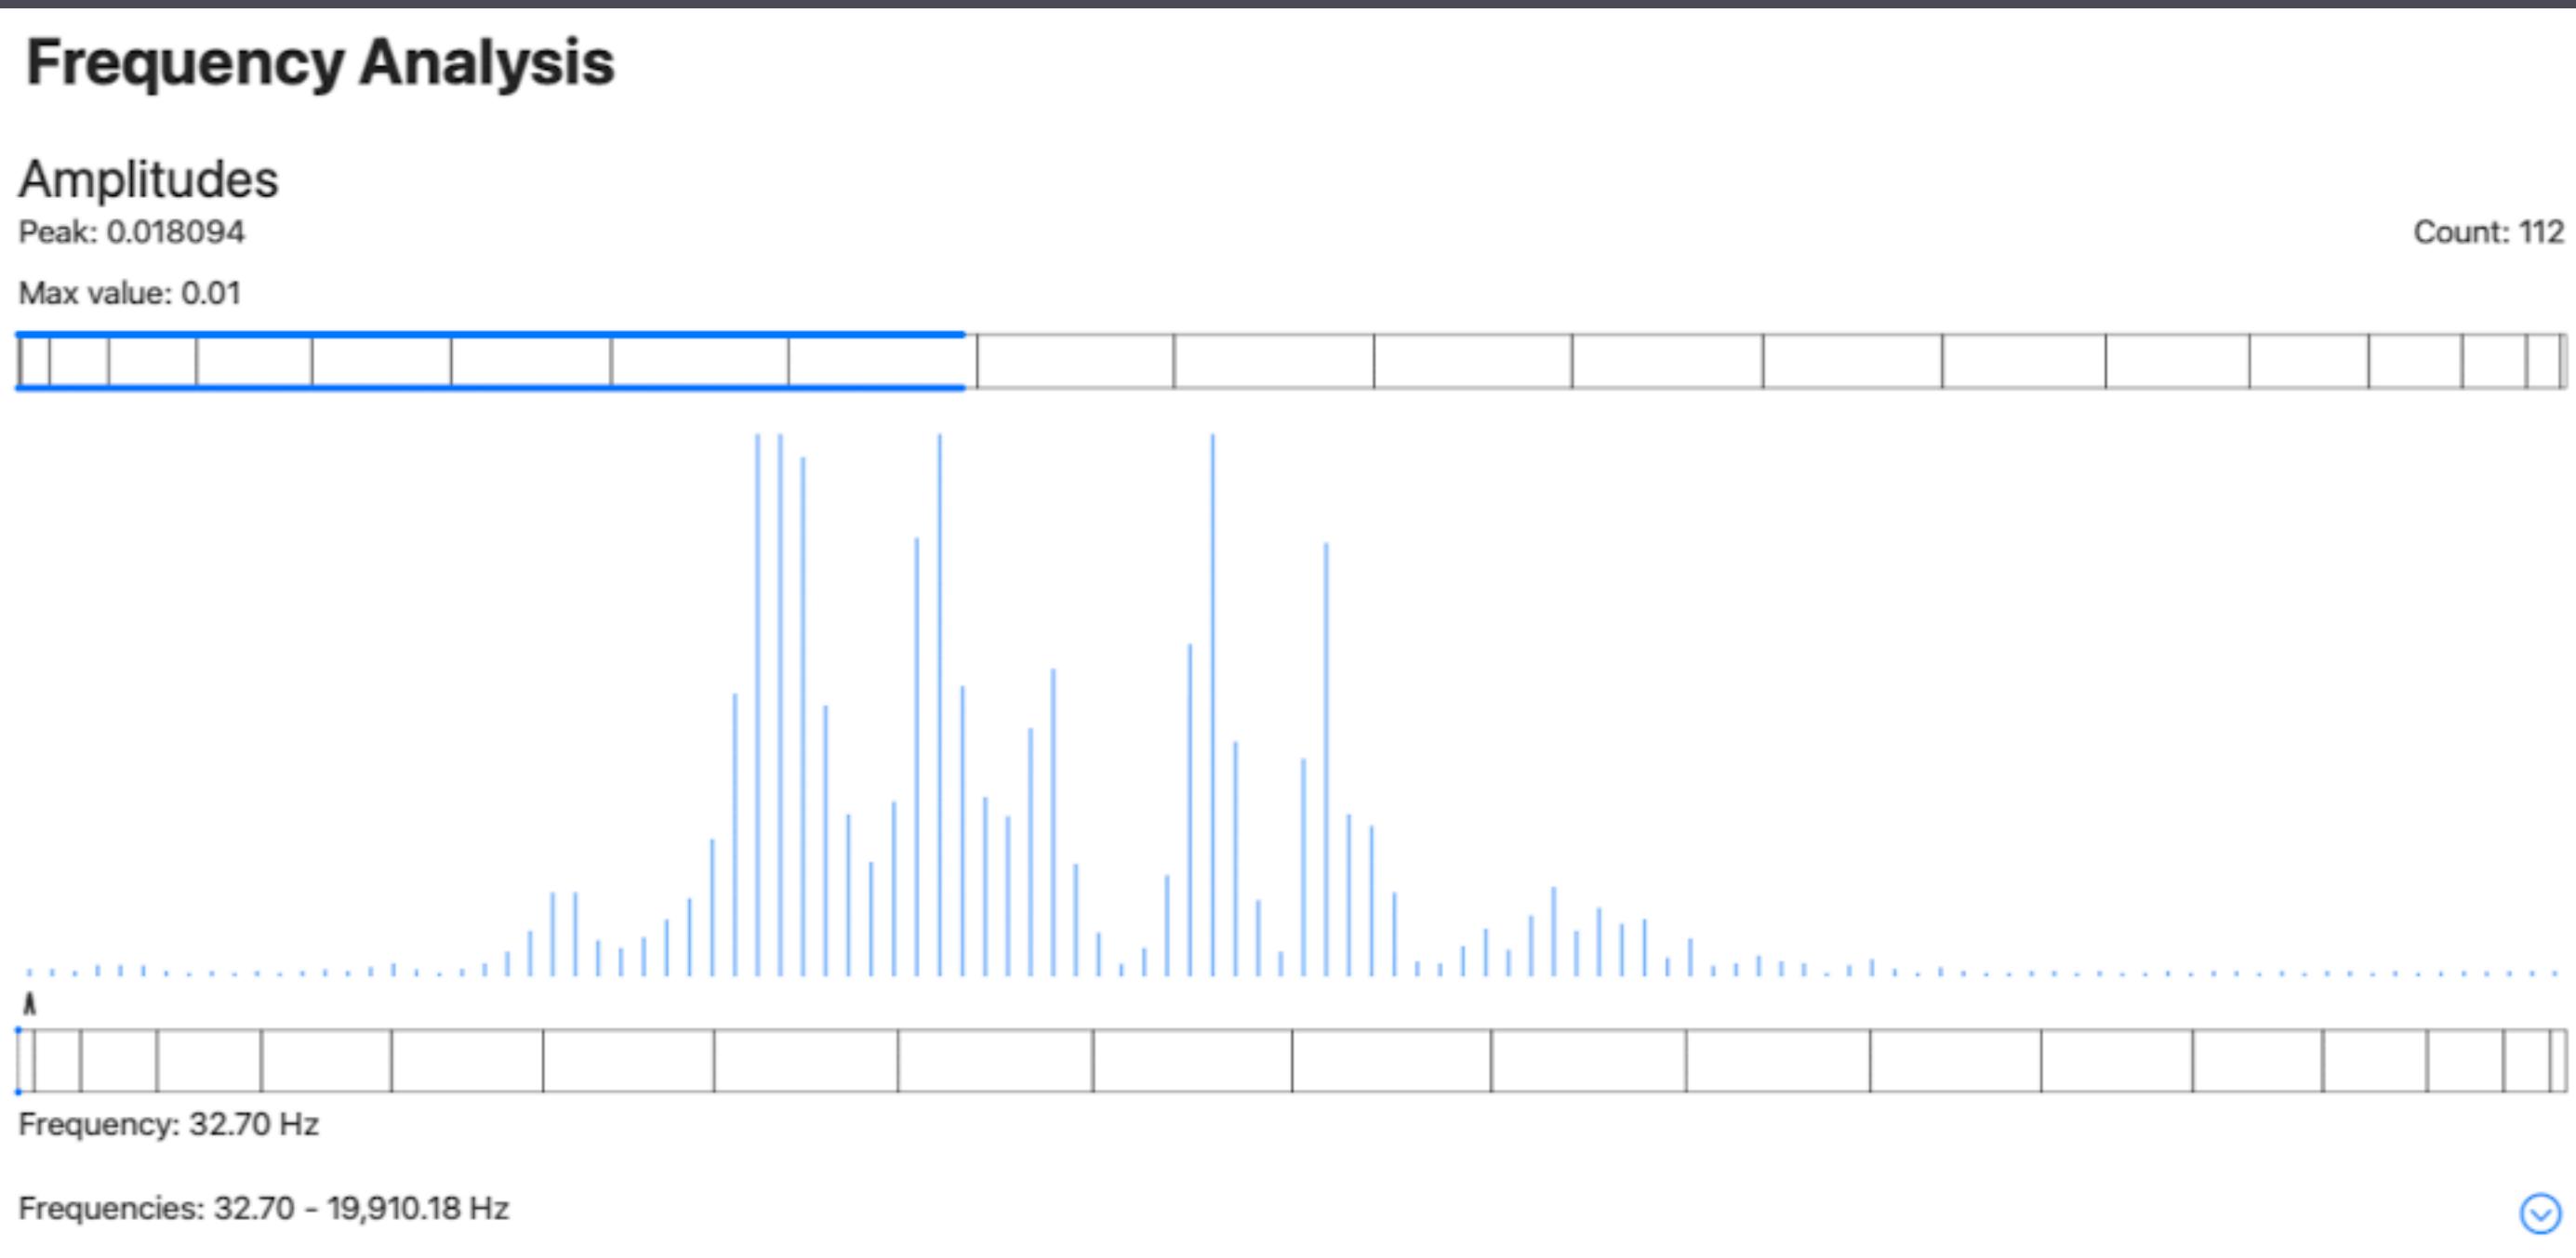


Outline

- Intro
- Single resonator
- **Resonator banks**
- Apps and applications

Resonator Banks

- Arbitrary number of resonators
- Directly tuned at frequencies of interest, arbitrarily distributed
 - Log frequency scale (like CQT), mel scale, etc.
 - Dynamics
- Resonators are independent
 - Parallel processing of each sample
- Analogous to bank of non-linear band pass filters



Resonator Banks

- Arbitrary number of resonators
- Directly tuned at frequencies of interest, arbitrarily distributed
 - Log frequency scale (like CQT), mel scale, etc.
 - Dynamics
- Resonators are independent
 - Parallel processing of each sample
- Analogous to bank of non-linear band pass filters

```
// Array of independent resonator instances
class ResonatorBankArray {
    var resonators = [Resonator]()
    [...]
}

// A bank of independent resonators
// implemented as a single array
// computations use the Accelerate framework
// with manual memory management
// (unsafe pointers)
class ResonatorBankVec {
    [...]

    // Accumulated resonance values
    // non-interlaced real (cos) | imaginary (sin) parts
    var rPtr : UnsafeMutableBufferPointer<Float>

    // Smoothed accumulated resonance values
    // non-interlaced real (cos) | imaginary (sin) parts
    var rrPtr : UnsafeMutableBufferPointer<Float>

    [...]
}
```

Resonator bank with Accelerate (SIMD)

```

// Update with sample                                // Single resonator
// compute alphas * sample                         // Update with sample
vDSP_vsmul(alphas, 1,                            alphaSample = alpha * sample
    &sample,
    alphasSample.baseAddress!, 1,
    vDSP_Length(twoNumResonators))

// resonator update
vDSP_vmma(rPtr.baseAddress!, 1,
          omAlphas, 1,
          zPtr.baseAddress!, 1,
          alphasSample.baseAddress!, 1,
          rPtr.baseAddress!, 1,
          vDSP_Length(twoNumResonators))

// smoothing using beta
vDSP_vmma(rrPtr.baseAddress!, 1,
          omBetas, 1,
          rPtr.baseAddress!, 1,
          betas, 1,
          rrPtr.baseAddress!, 1,
          vDSP_Length(twoNumResonators))

```

```

// Single resonator
// Update with sample
alphaSample = alpha * sample
omAlpha = 1.0 - alpha

c = omAlpha * c + alphaSample * Zc
s = omAlpha * s + alphaSample * Zs

omBeta = 1.0 - beta
cc = omBeta * cc + beta * c
ss = omBeta * ss + beta * s

```

Resonator bank with Accelerate (SIMD)

```

// Update with sample                                // all vectors of size 2 * numResonators
// compute alphas * sample                         // cosine and sine
vDSP_vsmul(alphas, 1,
            &sample,
            alphasSample.baseAddress!, 1,
            vDSP_Length(twoNumResonators))           // Advance phasor
                                                // Z <- Z * W
                                                // Z and W are DSPSplitComplex
vDSP_zvmul(&Z, 1,
            &W, 1,
            &Z, 1,
            vDSP_Length(numResonators),
            1)                                         // Re-normalization
                                                // vDSP.squareMagnitudes(Z, result: &smPtr)
vDSP_vmma(rPtr.baseAddress!, 1,
          omAlphas, 1,
          zPtr.baseAddress!, 1,
          alphasSample.baseAddress!, 1,
          rPtr.baseAddress!, 1,
          vDSP_Length(twoNumResonators))           // use reciprocal square root
                                                // vForce.rsqrt(smPtr, result: &rsqrtPtr)
                                                // vDSP.multiply(Z, by: rsqrtPtr, result: &Z)
// smoothing using beta
vDSP_vmma(rrPtr.baseAddress!, 1,
          omBetas, 1,
          rPtr.baseAddress!, 1,
          betas, 1,
          rrPtr.baseAddress!, 1,
          vDSP_Length(twoNumResonators))

```

Performance

Swift

Array of resonators

Frame Size	32 ◁
Implementation	Swift ◁
Update Model	Sequential ◁
Processing time per sample (ns):	19,736
Max samples per second:	50,670
Input Device	com.rogueamoeba.Loopb...82-A89B-99E57AC1FFFF ◁
Sample rate:	44,100 Hz

Frame Size	32 ◁
Implementation	Swift ◁
Update Model	Concurrent ◁
Processing time per sample (ns):	4,842
Max samples per second:	206,508
Input Device	com.rogueamoeba.Loopb...82-A89B-99E57AC1FFFF ◁
Sample rate:	44,100 Hz

Vectorized (SIMD)

Frame Size	32 ◁
Implementation	Vectorized ◁
Update Model	◊
Processing time per sample (ns):	396
Max samples per second:	2,526,249
Input Device	com.rogueamoeba.Loopb...82-A89B-99E57AC1FFFF ◁
Sample rate:	44,100 Hz

C++

Frame Size	32 ◁
Implementation	C++ ◁
Update Model	Sequential ◁
Processing time per sample (ns):	934
Max samples per second:	1,071,130
Input Device	com.rogueamoeba.Loopb...82-A89B-99E57AC1FFFF ◁
Sample rate:	44,100 Hz

Frame Size	32 ◁
Implementation	C++ ◁
Update Model	Concurrent ◁
Processing time per sample (ns):	1,404
Max samples per second:	712,425
Input Device	com.rogueamoeba.Loopb...82-A89B-99E57AC1FFFF ◁
Sample rate:	44,100 Hz

Frame Size	32 ◁
Implementation	Vectorized (C++) ◁
Update Model	◊
Processing time per sample (ns):	116
Max samples per second:	8,629,989
Input Device	com.rogueamoeba.Loopb...82-A89B-99E57AC1FFFF ◁
Sample rate:	44,100 Hz

Processing time per sample on MacBookPro M4 Pro - **debug mode!** (higher overhead for Swift versions)

Things NOT to do at Every Sample

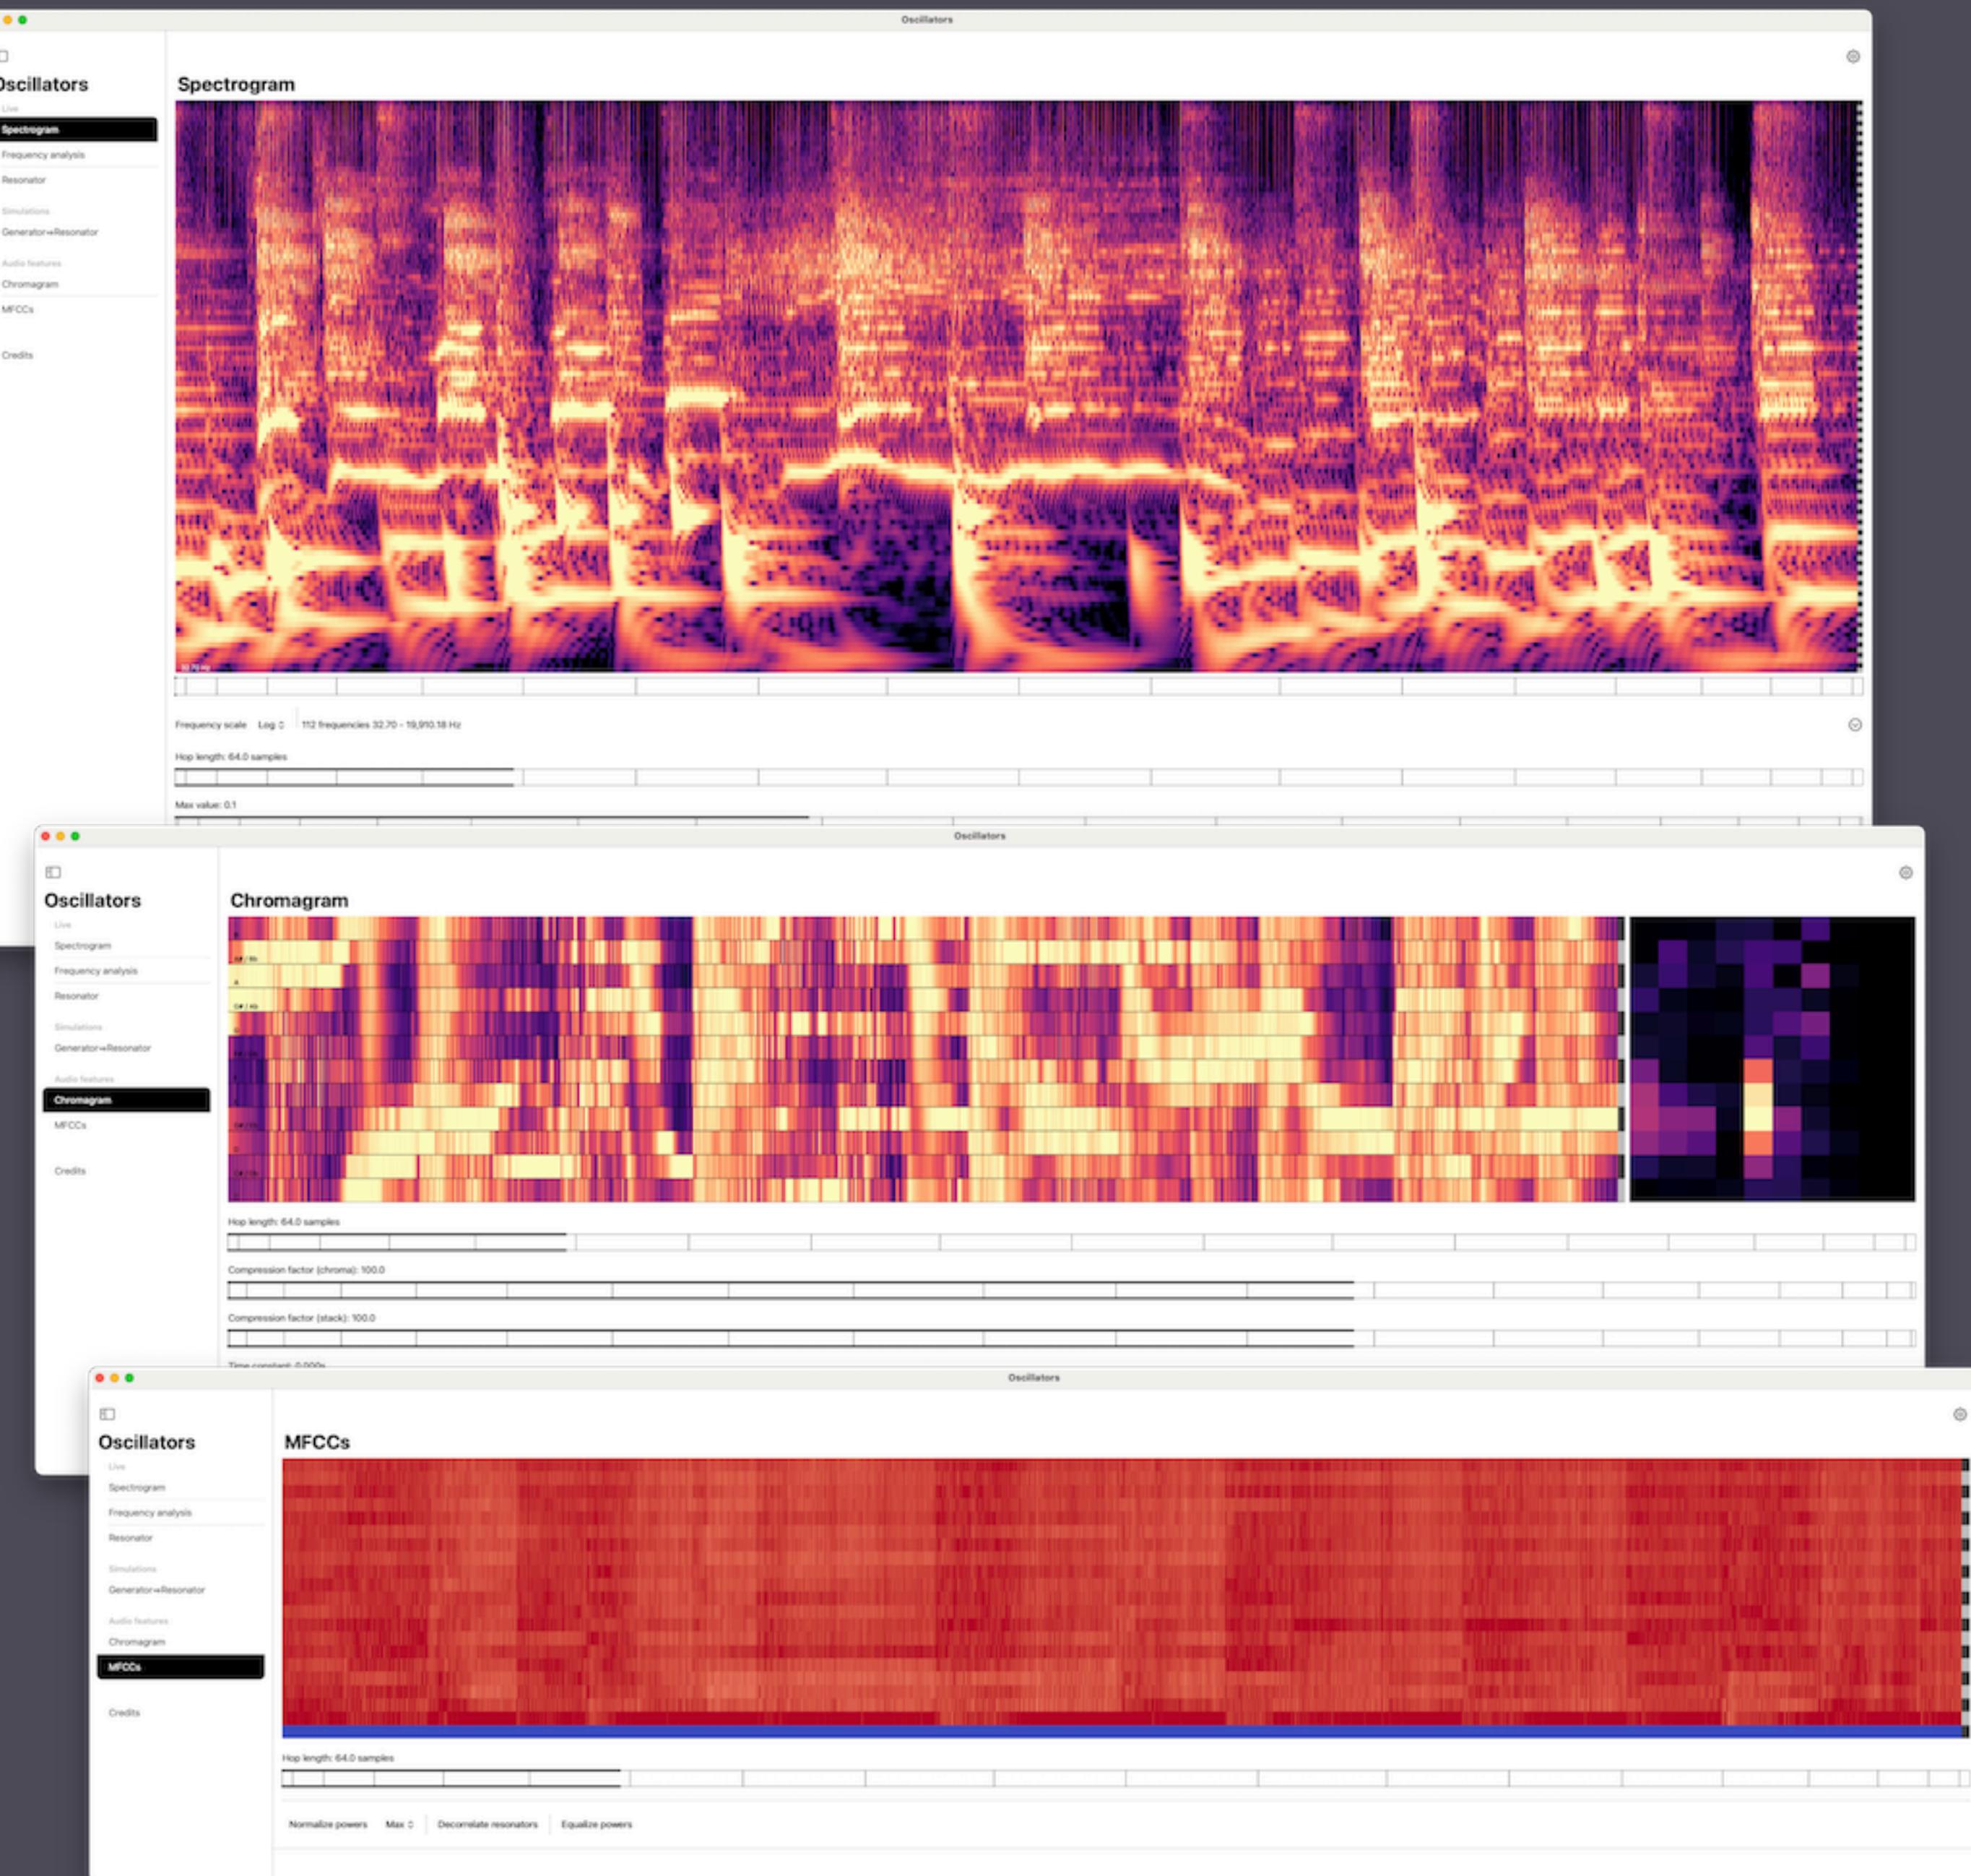
- Allocate buffer/arrays
 - Including those used in intermediary computations
- Copy/duplicate data
 - Most Accelerate functions can operate in-place (but not all...)
- Perform unnecessary computations
- Leak memory!

Outline

- Intro
- Single resonator
- Resonator banks
- **Apps and applications**

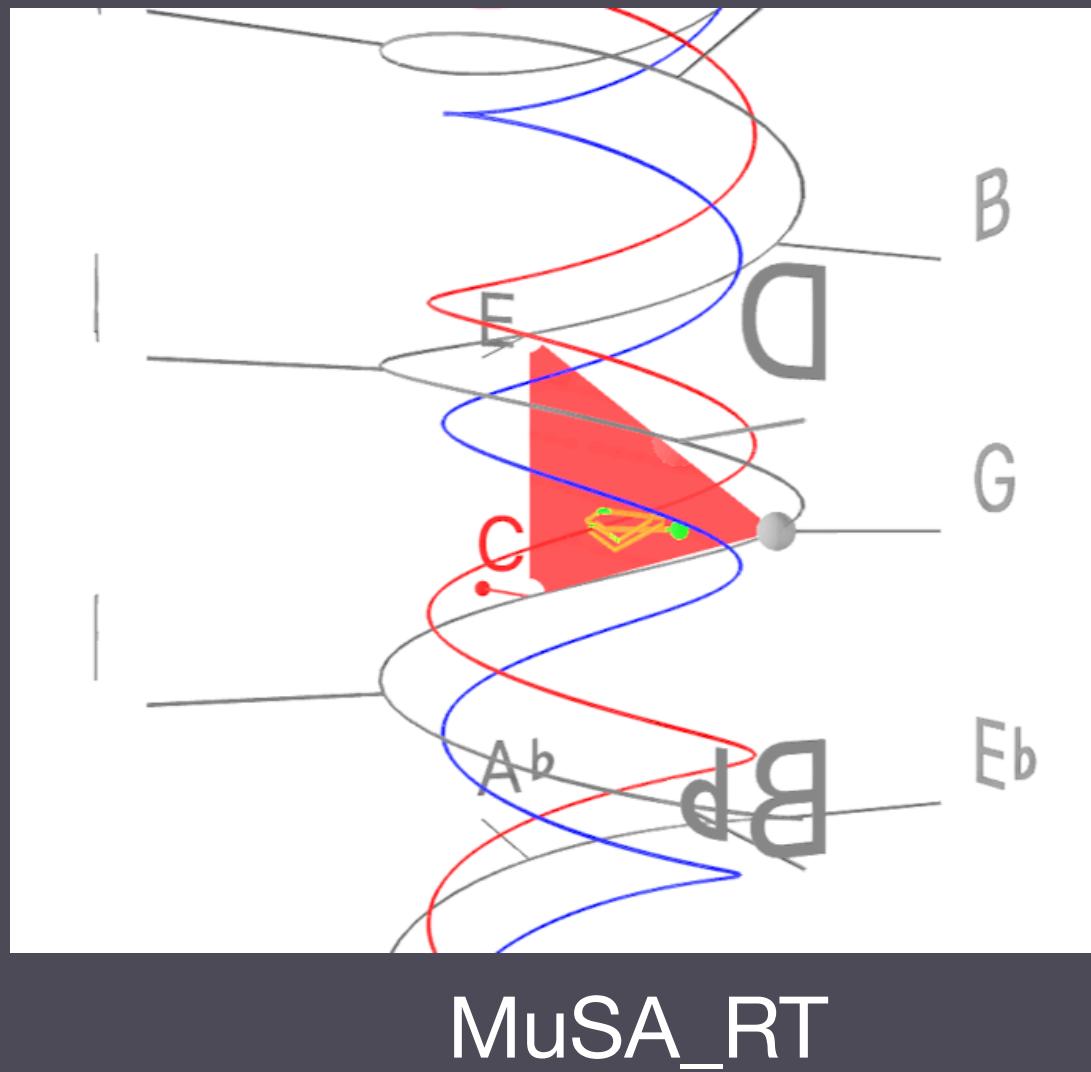
Oscillators App

- A playground for Oscillators Swift package
 - Single resonator
 - Resonator banks
 - **Spectrograms, chromagrams, MFCCs**
- AudioKit
- Swift UI
- Resonate Swift package
- Accelerate



Applications

- MuSA_RT?
- JUCE plugins?
- Optimized C++ implementation using Eigen for portability?
- ML
 - Replace Constant-Q transform: save memory and time; reduce inference latency (possibly not significantly compared to overall model latency, but it's a start...)
 - Efficient tokenization



MuSA_RT

Python wrapper (noFFT module)

- Resonate algorithm is auto regressive: cannot be parallelized across time
 - Usual Python vectorization does not work out of the box
 - Implement iteration loop in lower-level language (e.g. C++)
- noFFT module
 - C++ implementation with Python wrappers and utility functions
 - Jupyter notebooks
 - Uses Accelerate -> Eigen implementation for portability?
 - GPU implementation?

Thank you!



www.alexandrefrancois.org/Resonate